

**Yogesh Haribhau Kulkarni**

• You

...

AI Advisor (Helping organizations in their AI journeys) | PhD (Geometric Modeling) | Tech Colum...  
now • 

As an AI Advisor, I've been getting quite a few requests from companies, startups to large enterprises, exploring potential collaboration in GenAI and AI-driven initiatives.

The first step in any engagement is an exploratory call to understand where the company is at, on their AI journey, align on expectations, and identify real opportunities.

Here's a typical checklist I use to guide the conversation during that first call:

### Project Scope & Goals

- What are the immediate AI/ML project domains (next 6 to 12 months)?
- What longer-term AI opportunities or experiments are being considered?
- Any business KPIs or success metrics tied to these initiatives?

### Roadmap & Prioritization

- Do you have a 1 year AI roadmap? How fixed or flexible is it?
- Any prioritized use cases (e.g., chatbots, agents, automation, analytics)?
- What's driving prioritization: cost savings, efficiency, customer experience, etc.?

### Stakeholders & Teams

- Who are the internal stakeholders (AI/ML teams, PMs, DevOps, IT, etc.)?
- Are there clear owners for each initiative or deliverable?
- Who's responsible for governance, compliance, and risk management?

### Expectations & Reality Check

- Are you aware that GenAI is immature for mission-critical or non-text/image domains?
- Are current expectations aligned with realistic delivery (e.g., chatbots, co-pilots, basic agents only)?
- Is there awareness that LLM quality > agent design in most use cases?

### Tech Constraints & Risks

- Are you using or planning to use wrapper apps or black-box tools? (Risks: vendor lock-in, no debug-ability)
- What's your stance on open source vs. proprietary APIs (cost, control, performance)?
- Is there a plan to monitor model drift, data quality, and feedback loops?
- Is ML pipeline + human-in-loop (HITL) in place or being considered?

#### Infra, Cloud, and Scale

- Are workloads expected to run on cloud, hybrid, or on-prem?
- Do you have DevOps/MLOps maturity for scaling and productionizing?
- How is security, observability, and compliance handled in production workflows?


#### Data & Foundation Models

- Do you have clean, labeled data for fine-tuning or evaluation?
- Any existing foundation models or vendors in use?
- Are you planning on building custom models or using off-the-shelf?

#### Integration & Tooling

- What is your existing tech stack (data, APIs, CI/CD, monitoring)?
- Any plans for agent orchestration or chaining tools (LangChain, Haystack, etc.)?
- Preferred frameworks, languages, or platforms?

Anything missing? Let me know in the comments...

 If you're considering GenAI adoption and want to run a sanity check, happy to talk. More details about my own AI Advisory approach (including commercials) is at <https://lnkd.in/dhqWbeHS>

[#AI](#) [#GenAI](#) [#MLOps](#) [#EnterpriseAI](#) [#StartupGrowth](#) [#AIAdoption](#) [#LLM](#) [#ChatGPT](#) [#Productivity](#)  
[#TechLeadership](#) [#ArtificialIntelligence](#) [#FutureOfWork](#)

