

**Yogesh Haribhau Kulkarni**

• You

AI Advisor (Helping organizations in their AI journeys) | PhD (Geometric Modeling) | Tech Colum...
now •

...

Building an Intelligent FAQ Chatbot with RAG!

Ever noticed how traditional FAQ (Frequently-Asked-Questions) systems force users to guess exact keywords? 🤔 A customer asking "How long for delivery?" gets no match when your FAQ says "What is the shipping time?", same question, zero results. This isn't just frustrating; it's costing businesses thousands in support tickets.

Here is an example RAG-based FAQ chatbot that thinks like humans, not search engines, at [GitHub](#)
<https://lnkd.in/dF-jDWVs>

It has Metadata-First Architecture ie Questions vectorized for matching, answers stored as metadata. This prevents answer content from polluting semantic search.

The result? Semantic understanding with good accuracy, fast responses, and it runs entirely on open-source models ([Hugging Face](#)). No expensive API lock-ins. 💰

Real-world impact: Banking FAQ with 1000+ entries now handles natural language variations seamlessly.

Read the full technical deep-dive at [Medium](#) <https://lnkd.in/dM2rcaHd>

[#MachineLearning](#) [#RAG](#) [#AI](#) [#Chatbot](#) [#NLP](#) [#LlamaIndex](#) [#VectorSearch](#) [#CustomerExperience](#)
[#SemanticSearch](#) [#GenAI](#) [#TechInnovation](#) [#Python](#) [#OpenSource](#)

