

**Yogesh Haribhau Kulkarni**

• You

...

AI Advisor (Helping organizations in their AI journeys) | PhD (Geometric Modeling) | Tech Colum...
now •

Almost everyday I get LinkedIn feed having posts like "Top 10 research papers you must read to become an AI expert." 📖 Huuuushhhh... 😩

It is not that I do not want to study them well. It is that there are only 24 hours in a day, and reading every paper the "Deep" way feels impossible.

Traditionally, there are two approaches:

- ✅ Quick way: Skim the abstract, check the figures and tables, read the proposal and conclusions.
- ✅ Deep way: Read every line, take notes, and study the details.

But then, now there is a third way, the AI way 🤖

With Notebook LM by [Google](#), you just need to:

- 1 Upload the paper, that's it.
- 2 You get a concise summary
- 3 Chat with it to clarify concepts
- 4 Generate a 15–17 minute podcast 🎙️
- 5 Create a short video explainer 🎥
- 6 Get a ready visual mind map 🧠

That podcast is pure magic. You have to try it to understand.

I tested it on the recent Kimi K1.5 paper, a remarkable new LLM. Below are some of those pics.

Well, I understand that this will never fully replace the "deep way", but as a filter and a fast track to understanding, it is a game-changer. 🚀

Do try and let me know

#ArtificialIntelligence #MachineLearning #LLM #GoogleAI #Research #Productivity #Altools

#DeepLearning #Innovation #TechTips

#NotebookLM #Altools #Notetaking #TreeDiagram #SmartResearch #Productivity #StudentTools

#GoogleWorkspace #FutureOfWork #VisualLearning #MindMapping #DeepThinking

Kimi K1.5: Scaling RL with LLMs

Share

Settings

Sources

+ Add

Discover

Select all sources

Kimi Scaling Reinforcement Learning with LLMs.pdf

Chat

Kimi K1.5: Scaling RL with LLMs

1 source

The provided source details the Kimi K1.5, a multi-modal large language model (LLM) primarily trained using reinforcement learning (RL). It introduces various innovative techniques such as long context scaling, improved policy optimization, and partial rollouts to enhance its reasoning capabilities, particularly for generating long chains of thought (CoT). The report also describes methods for improving short-CoT models and highlights the infrastructure optimizations that support large-scale RL training. Ultimately, the document showcases Kimi K1.5's state-of-the-art performance across numerous benchmarks in text, reasoning, and vision, demonstrating the effectiveness of RL in scaling LLM intelligence.

Save to note

Add note

Audio Overview

Mind Map

Start typing...

1 source

How do long context scaling and improved policy optimization enhance LLM reasoning capabilities? What are the key

Studio

Audio Overview

Video Overview

Mind Map

Reports

Studio output will be saved here.

After adding sources, click to add Audio Overview, Study Guide, Mind Map, and more!

Add note

Kimi K1.5: Scaling RL with LLMs

Training Pipeline Stages

Reinforcement Learning Details

LongShort: Context Compression for Short-CoT Models

Other Training Details

RL Infrastructure

Problem Setting

Policy Optimization

Length Penalty

Sampling Strategies

More Details on Training Recipe

Transfers thinking prices from long-CoT to short-CoT

Improves performance with limited test-time tokens

Approaches

Kimi K1.5: Scaling RL with LLMs

Sources

+ Add

Discover

Select all sources

Kimi Scaling Reinforcement Learning with LLMs.pdf

Chat

Kimi K1.5: Scaling RL with LLMs

1 source

The provided source refers to "DeepSeek V3" rather than "Deepseek V3". Assuming your query refers to Deepseek V3 as presented in the source, here is a brief overview of the model's architecture and training details. **Architecture:** Deepseek V3 is a multi-modal large language model (LLM) primarily trained on both text and vision data, designed to reason across these modalities. It features a unified architecture, enabling cross-modal reasoning and multimodal understanding. **Training Details:** The model is trained using a combination of supervised fine-tuning (SFT) and reinforcement learning (RL). It leverages a large-scale dataset and a sophisticated training infrastructure to achieve state-of-the-art performance. **Performance:** Deepseek V3 demonstrates superior performance across various benchmarks, particularly in long context scaling and reasoning tasks. It outperforms many existing models in these domains, showcasing its ability to handle complex tasks and generate high-quality outputs. **Key Features:** The model is designed to be efficient and scalable, making it suitable for a wide range of applications. It supports a variety of input modalities and can generate outputs in multiple formats, including text, images, and audio. **Conclusion:** Deepseek V3 represents a significant advancement in the field of large language models, offering a powerful and versatile tool for a variety of tasks. Its performance and capabilities make it a strong contender for future research and development in the field of AI.

Save to note

Audio Overview

Mind Map

Start typing...

1 source

How do long context scaling and improved policy optimization enhance LLM reasoning capabilities? What are the key

Studio

Audio Overview

Video Overview

Mind Map

Reports

Studio output will be saved here.

After adding sources, click to add Audio Overview, Study Guide, Mind Map, and more!

Add note

https://www.linkedin.com/feed/update/urn:li:ugcPost:7359473715048341504/

2/2