

INTRODUCTION TO LANGCHAIN

Yogesh Haribhau Kulkarni



Outline

① INTRODUCTION

② FRAMEWORK

③ WHAT'S NEW

④ CONCLUSIONS

⑤ REFERENCES

Introduction to LangChain

YHK

What is LangChain?

LangChain: A comprehensive framework for building LLM-powered applications

- ▶ **Core Purpose:** Simplify development of applications using LLMs
- ▶ **Key Solutions:**
 - ▶ **RAG:** Connect language models to external data sources
 - ▶ **Agentic:** Allow language models to interact with their environment
- ▶ **Main Features:**
 - ▶ Generic interface to various foundation models
 - ▶ Advanced prompt management framework
- ▶ **Availability:** Python and JavaScript libraries
- ▶ **Open Source:** MIT License, created by Harrison Chase
- ▶ **Repository:** <https://github.com/langchain-ai/langchain>

(Ref: Getting Started with LangChain: A Beginner's Guide to Building LLM-Powered Applications)

Why You Need LangChain

- ▶ LLMs alone lack context, memory, and retrieval abilities.
- ▶ Chatbots need to manage chat history and company knowledge bases.
- ▶ Storing, retrieving, and reasoning over data is complex manually.
- ▶ LangChain acts as an abstraction layer for building AI agents.
- ▶ It integrates LLMs, memory, and tools seamlessly.
- ▶ Enables vendor flexibility, easy switch between Open/Close models.
- ▶ Reduces code complexity and accelerates AI app development.
- ▶ Simplifies connecting models, databases, and APIs into a single framework.

LLMs vs Agentic Software

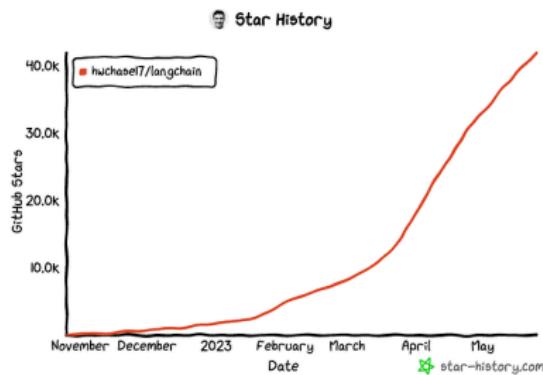
- ▶ LLMs are static, respond from training data without awareness.
- ▶ Agents have autonomy, memory, and tool use for dynamic actions.
- ▶ Example: Refund query, agent retrieves policy, product, and chat history.
- ▶ Agents can access vector databases and retrieve company knowledge.
- ▶ LangChain enables memory persistence across user conversations.
- ▶ Traditional software follows fixed logic; agents make adaptive decisions.
- ▶ Agentic software uses modular components (LLMs, memory, retrievers).
- ▶ LangChain provides prebuilt tools for these agentic capabilities.

LangChain Core Components and Labs

- ▶ Core modules: LLM connectors, memory, embeddings, vector stores.
- ▶ Simplifies LLM API setup, e.g., ChatOpenAI() or ChatGroq()
- ▶ Supports databases like Chroma or Pinecone for knowledge retrieval.
- ▶ Prompt templates manage dynamic, reusable prompts and chat context.
- ▶ LCEL (LangChain Expression Language) composes chains without legacy LLMChain wrappers
- ▶ Enables async, streaming, and batch workflows with type safety.

So, Why LangChain?

- ▶ **RAG Applications:** Build Retrieval Augmented Generation apps with external data
- ▶ **Agent Systems:** Create intelligent agents with tool access
- ▶ **Production Ready:** Logging, callbacks, monitoring built-in
- ▶ **Model Agnostic:** Switch between LLM providers easily
- ▶ **Modular Design:** Compose components as needed



(Ref: LangChain tutorial: Build an LLM-powered app)

LangChain Community & Ecosystem

- ▶ **Community Growth:**
 - ▶ Over 80,000+ GitHub stars (as of 2024)
 - ▶ 2,000+ contributors
 - ▶ Millions of monthly downloads
 - ▶ Active Discord, Twitter/X presence
- ▶ **License:** MIT License - freely modifiable and commercial-friendly
- ▶ **Ecosystem Components:**
 - ▶ **LangChain Core:** Foundation library
 - ▶ **LangGraph:** Stateful, multi-actor applications
 - ▶ **LangServe:** Deploy as REST APIs
 - ▶ **LangSmith:** Debugging, testing, monitoring

(Ref: What is Langchain and why should I care as a developer? - Logan Kilpatrick)

Installation & Setup

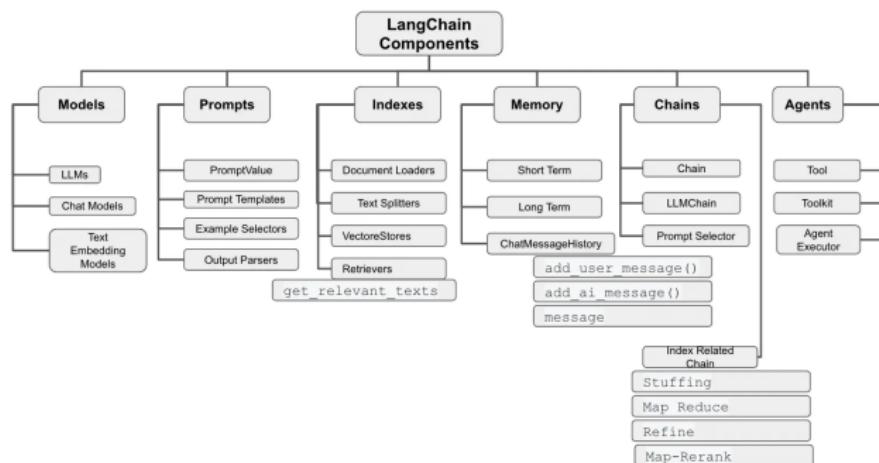
Prerequisites:

- ▶ Python version ≥ 3.9 and < 4.0
 - ▶ Modern Installation (2024)
 - ▶ (Ref: Official LangChain Documentation 2024)
- :

```
1 # Core packages
2 pip install langchain langchain-core
3
4 # Provider-specific packages for Groq
5 pip install langchain-groq
6
7 # Alternative providers
8 # pip install langchain-anthropic langchain-google-genai
9
10 # Community integrations for embeddings
11 pip install langchain-community langchain-huggingface sentence-transformers
12
13 # Optional components
14 pip install chromadb          # Vector store
15 pip install langchain-text-splitters # Document processing
16
17 //API Keys Setup:
18
19 import os
20 os.environ["GROQ_API_KEY"] = "your-groq-api-key-here"
21 # Or use .env file with python-dotenv
```



Core Components Overview



- ▶ **Models:** LLMs, Chat Models, Embeddings
- ▶ **Prompts:** Template management and optimization
- ▶ **Chains/LCEL:** Compose components with pipes
- ▶ **Memory:** Conversation state management
- ▶ **Retrievers:** Access external data
- ▶ **Agents:** Dynamic tool selection and execution

Modern LangChain: LCEL (LangChain Expression Language)

What is LCEL?

- ▶ Modern, declarative way to build chains (introduced 2023)
- ▶ Uses pipe operator | to chain components
- ▶ Replaces old LLMChain pattern
- ▶ Built-in streaming, async, and batch support

Key Benefits:

- ▶ **Simplicity:** More readable and concise
- ▶ **Streaming:** Built-in streaming by default
- ▶ **Async:** Native async/await support
- ▶ **Observability:** Better tracing and debugging
- ▶ **Fallbacks:** Easy error handling and retries

(Ref: LangChain LCEL Documentation 2024)

LCEL: Basic Example

```
1 // Old Pattern (Deprecated):
2 from langchain.llms import OpenAI
3 from langchain.chains import LLMChain
4
5 llm = OpenAI()
6 chain = LLMChain(llm=llm, prompt=prompt)
7 result = chain.run("input")
8
9 // Modern LCEL Pattern with Groq:
10
11 from langchain.groq import ChatGroq
12 from langchain.core.prompts import ChatPromptTemplate
13 from langchain.core.output_parsers import StrOutputParser
14
15 # Set up Groq model, e.g., Gemma or Llama 3
16 llm = ChatGroq(model.name="gemma-7b-it")
17 prompt = ChatPromptTemplate.from_template("Tell me about {topic}")
18 output_parser = StrOutputParser()
19
20 # Build chain with pipe operator
21 chain = prompt | llm | output_parser
22
23 # Invoke the chain
24 result = chain.invoke({ "topic": "LangChain" })
```

LCEL: Advanced Features

```
// Streaming Support:  
2 chain = prompt | llm | output_parser  
4  
# Stream tokens as they're generated  
6 for chunk in chain.stream({"topic": "AI"}):  
    print(chunk, end="", flush=True)  
8  
// Async Execution:  
10 # Async invocation  
result = await chain.ainvoke({"topic": "AI"})  
12  
# Async streaming  
14 async for chunk in chain.astream({"topic": "AI"}):  
    print(chunk, end="", flush=True)  
16  
// Batch Processing:  
18 # Process multiple inputs in parallel  
results = chain.batch([  
    {"topic": "AI"},  
    {"topic": "ML"},  
    {"topic": "LangChain"}  
])
```

LCEL: Complex Chains

```
1 // Multi-step Chain with RunnablePassthrough:  
from langchain.core.runnables import RunnablePassthrough  
3  
chain = (  
    {"context": retriever, "question": RunnablePassthrough()  
    | prompt  
    | llm  
    | output_parser  
)  
11 result = chain.invoke("What is LangChain?")  
  
13 // Parallel Execution with RunnableParallel:  
  
15 from langchain.core.runnables import RunnableParallel  
  
17 chain = RunnableParallel(  
    summary=prompt1 | llm | output_parser,  
    keywords=prompt2 | llm | output_parser  
)  
21  
23 result = chain.invoke( { "text": "Long document..." } )  
# Returns: { "summary": "...", "keywords": "..." }
```

Model Integration: Modern Approach

```
1 Updated Import Structure:  
3 # OpenAI models  
4 from langchain.openai import ChatOpenAI, OpenAIEmbeddings  
5  
# Hugging Face models  
6 from langchain.huggingface import HuggingFaceEndpoint  
7  
# Google models (Gemini)  
8 from langchain.google.genai import ChatGoogleGenerativeAI  
11  
# Anthropic models  
12 from langchain.anthropic import ChatAnthropic  
13  
// Example Usage:  
14 from langchain.groq import ChatGroq  
15 from langchain.core.messages import HumanMessage, SystemMessage  
16  
# Initialize Groq with Gemma  
17 chat = ChatGroq(model="gemma2-9b-it", temperature=0.7)  
18  
messages = [  
19     SystemMessage(content="You are a helpful assistant"),  
20     HumanMessage(content="Explain LangChain in 2 sentences")  
21 ]  
22  
23 response = chat.invoke(messages)  
24 print(response.content)
```

Complete RAG Application: A Quick Start Example

```
from langchain.groq import ChatGroq
2 from langchain.community.embeddings import HuggingFaceEmbeddings
from langchain.community.document.loaders import WebBaseLoader
4 from langchain.textsplitters import RecursiveCharacterTextSplitter
from langchain.community.vectorstores import Chroma
6 from langchain.core.prompts import ChatPromptTemplate
from langchain.core.output_parsers import StrOutputParser
8 from langchain.core.runnables import RunnablePassthrough

10 # Load and process documents
loader = WebBaseLoader("https://example.com/doc")
12 docs = loader.load()
splitter = RecursiveCharacterTextSplitter(chunk_size=1000)
14 splits = splitter.split_documents(docs)

16 embeddings = HuggingFaceEmbeddings()
vectorstore = Chroma.from_documents(splits, embeddings)
18 retriever = vectorstore.as_retriever()

20 prompt = ChatPromptTemplate.from_template("""
Answer based on context: {context}
22 Question: {question}""")
24 model = ChatGroq(model.name="llama3-8b-8192")

26 # RunnablePassthrough => "context": retriever("What is the main topic?"),
chain = (
28 {"context": retriever, "question": RunnablePassthrough()}
| prompt | model | StrOutputParser())
30 response = chain.invoke("What is the main topic?")
```

Key Concepts: Chains vs Agents

- ▶ **Chains (LCEL):**
 - ▶ Predetermined sequence of operations
 - ▶ Composed with pipe operator: `prompt | llm | parser`
 - ▶ Fixed execution path
 - ▶ Best for: Structured, predictable workflows
- ▶ **Agents:**
 - ▶ Dynamic decision-making with LLM reasoning
 - ▶ Choose tools based on input
 - ▶ Adaptive execution path
 - ▶ Best for: Complex, unpredictable scenarios

When to use what?

- ▶ Use **Chains/LCEL** when: Steps are known, workflow is fixed
- ▶ Use **Agents** when: Need dynamic tool selection, multi-step reasoning

(Ref: Superpower LLMs with Conversational Agents)

LangChain Use Cases

- ▶ **Retrieval Augmented Generation (RAG):**
 - ▶ Document Q&A systems
 - ▶ Knowledge base search
 - ▶ Semantic search applications
- ▶ **Conversational AI:**
 - ▶ Chatbots with memory
 - ▶ Customer support agents
 - ▶ Personal assistants
- ▶ **Data Analysis:**
 - ▶ SQL query generation
 - ▶ Tabular data Q&A
 - ▶ Report generation
- ▶ **Autonomous Agents:**
 - ▶ Web scraping and research
 - ▶ Multi-tool workflows
 - ▶ API integration

(Ref: LangChain Use Cases Documentation)

LangChain Framework Components

Framework Architecture

Chains: The core of LangChain. Components (and even other chains) can be stringed together to create *chains*.

Prompt templates: Prompt templates are templates for different types of prompts. Like “chatbot” style templates, ELI5 question-answering, etc

LLMs: Large language models like GPT-3, BLOOM, etc

Indexing Utils: Ways to interact with specific data (embeddings, vectorstores, document loaders)

Tools: Ways to interact with the outside world (search, calculators, etc)

Agents: Agents use LLMs to decide what actions should be taken. Tools like web search or calculators can be used, and all are packaged into a logical loop of operations.

Memory: Short-term memory, long-term memory.

Core Building Blocks:

- ▶ **Models:** LLMs, Chat Models, Embeddings
- ▶ **Prompts:** Dynamic template management
- ▶ **Output Transformers:** Structured output extraction
- ▶ **Retrievers:** Document and data access
- ▶ **Memory:** Conversation state persistence
- ▶ **Agents & Tools:** Dynamic reasoning and actions

(Ref: Building the Future with LLMs, LangChain, & Pinecone)



Models

YHK

Models in LangChain

Three Types of Models:

- ▶ **LLMs (Large Language Models):**
 - ▶ Input: String (prompt)
 - ▶ Output: String (completion)
 - ▶ Examples: GPT-4, Claude, Gemma, Llama 3, Mixtral.
 - ▶ Use case: Text generation, completion
- ▶ **Chat Models:**
 - ▶ Input: List of messages
 - ▶ Output: Chat message
 - ▶ Examples: ChatGPT, Claude Chat, ChatGroq
 - ▶ Use case: Conversational AI
- ▶ **Embedding Models:**
 - ▶ Input: Text
 - ▶ Output: Vector (list of floats)
 - ▶ Examples: OpenAI Embeddings, HuggingFace, Sentence Transformers
 - ▶ Use case: Semantic search, similarity

Model Integration: Modern Syntax

```
1 from langchain.groq import ChatGroq
2 from langchain.community.embeddings import HuggingFaceEmbeddings
3 from langchain.core.messages import HumanMessage, SystemMessage
4 from langchain.core.output_parsers import StrOutputParser
5
6 # Chat Model with Groq
7 chat = ChatGroq(model.name="llama3-8b-8192", temperature=0.7)
8
9 messages = [
10     SystemMessage(content="You are a helpful assistant"),
11     HumanMessage(content="Explain quantum computing briefly")
12 ]
13 response = chat.invoke(messages)
14
15 # Embedding Model (from Hugging Face)
16 embeddings = HuggingFaceEmbeddings(model_name="all-MiniLM-L6-v2")
17 vector = embeddings.embed_query("Machine learning is...")
18
19 # Using with LCEL
20 from langchain.core.prompts import ChatPromptTemplate
21
22 prompt = ChatPromptTemplate.from_template("Explain {topic}")
23 chain = prompt | chat | StrOutputParser()
24 result = chain.invoke({ "topic": "blockchain" })
```

Alternative Model Providers

```
// Groq (Recommended for Fast Inference):
2 from langchain.groq import ChatGroq

4 # Gemma 2 — Very fast, efficient
Ilm = ChatGroq(model.name="gemma2-9b-it", temperature=0.7)
6 response = Ilm.invoke("What is the Groq LPU?")

8 # Llama 3 — Balanced performance
Ilm = ChatGroq(model.name="llama3-70b-8192", temperature=0.7)

10 // Hugging Face (Self-hosted):
12 from langchain.huggingface import HuggingFaceEndpoint
Ilm = HuggingFaceEndpoint(
14     repo_id="mistralai/Mistral-7B-Instruct-v0.2",
15     temperature=0.7
16 )

18 // Google Gemini (Multimodal):
from langchain.google.genai import ChatGoogleGenerativeAI
20 Ilm = ChatGoogleGenerativeAI(model="gemini-pro", temperature=0.7)

22 // Anthropic Claude (Reasoning):
from langchain.anthropic import ChatAnthropic
24 Ilm = ChatAnthropic(model="claude-3-5-sonnet-20241022")
```

Prompts

YHK

Prompts in LangChain

Prompt Management Strategies:

- ▶ **Prompt Templates:**
 - ▶ Parameterized templates with variables
 - ▶ Dynamic input insertion
 - ▶ Reusable prompt structures
- ▶ **Chat Prompt Templates:**
 - ▶ Multi-message conversations
 - ▶ System, human, AI message roles
 - ▶ Better for chat models
- ▶ **Few-Shot Prompts:**
 - ▶ Include example inputs/outputs
 - ▶ Guide model response style
 - ▶ Improve accuracy

Prompt Templates: Modern Examples

Basic Prompt Template:

```
from langchain.core.prompts import PromptTemplate
2
template = """You are a {role} assistant.
4 Task: {task}
Context: {context}
6 Provide a {format} response."""
8 prompt = PromptTemplate(
    template=template,
10    input_variables=["role", "task", "context", "format"]
)
12
formatted = prompt.format(
14    role="helpful",
    task="explain quantum computing",
16    context="for beginners",
    format="simple"
18 )
20 // Chat Prompt Template:
from langchain.core.prompts import ChatPromptTemplate
22
prompt = ChatPromptTemplate.from_messages([
24    ("system", "You are a {role}"),
    ("human", "{input}"),
26    ("ai", "I understand. Let me help with that."),
    ("human", "{follow-up}")
28 ])
```

Few-Shot Prompting

```
from langchain.core.prompts import FewShotPromptTemplate, PromptTemplate
2
# Define examples
4 examples = [
5     {"input": "happy", "output": "joyful, cheerful, delighted"}, 
6     {"input": "sad", "output": "unhappy, sorrowful, dejected"}, 
7 ]
8
# Example template
10 example_prompt = PromptTemplate(
11     input_variables=["input", "output"],
12     template="Input: {input}\nSynonyms: {output}"
13 )
14
# Few-shot template
16 few_shot_prompt = FewShotPromptTemplate(
17     examples=examples,
18     example_prompt=example_prompt,
19     prefix="Provide synonyms for the following word:",
20     suffix="Input: {word}\nSynonyms:",
21     input_variables=["word"]
22 )
23
# Use in chain with Groq
24 from langchain.groq import ChatGroq
25 from langchain.core.output_parsers import StrOutputParser
26
27 chain = few_shot_prompt | ChatGroq(model.name="gemma-7b-it") | StrOutputParser()
28 result = chain.invoke({ "word": "angry" })
```



Memory

YHK

Memory in LangChain

Why Memory?

- ▶ LLMs are stateless by default
- ▶ Chatbots need conversation context
- ▶ Memory stores and retrieves conversation history

Memory Types (Consolidated):

- ▶ **ConversationBufferMemory:**
 - ▶ Stores entire conversation history
 - ▶ Simple but can exceed token limits
 - ▶ Best for: Short conversations
- ▶ **ConversationBufferWindowMemory:**
 - ▶ Keeps only last K interactions
 - ▶ Prevents token overflow
 - ▶ Best for: Longer conversations with recent context
- ▶ **ConversationSummaryMemory:**
 - ▶ Summarizes old messages
 - ▶ Reduces token usage
 - ▶ Best for: Very long conversations

Memory: Implementation Examples

Buffer Memory (Stores Everything):

```
1 from langchain.memory import ConversationBufferMemory
2 from langchain_groq import ChatGroq
3 from langchain_core.prompts import ChatPromptTemplate, MessagesPlaceholder
4 from langchain_core.runnables import RunnablePassthrough
5
6 memory = ConversationBufferMemory(
7     return_messages=True,
8     memory_key="history"
9 )
10
11 prompt = ChatPromptTemplate.from_messages([
12     ("system", "You are a helpful assistant"),
13     MessagesPlaceholder(variable_name="history"),
14     ("human", "{input}")
15 ])
16
17 chain = (
18     RunnablePassthrough.assign(
19         history=lambda x: memory.load_memory_variables({})[{"history"}]
20     )
21     | prompt
22     | ChatGroq(model_name="gemma2-9b-it")
23 )
24
25 # Use the chain
26 response = chain.invoke({"input": "Hi, I'm Alice"})
27 memory.save_context({"input": "Hi, I'm Alice"}, {"output": response.content})
```

Memory: Window Memory

Benefits:

- ▶ Fixed memory footprint
- ▶ Maintains recent context
- ▶ Prevents token limit issues

Buffer Window Memory (Keeps Last K):

```
1 from langchain.memory import ConversationBufferWindowMemory
2
3 memory = ConversationBufferWindowMemory(
4     k=3, # Only keeps last 3 interactions
5     return_messages=True,
6     memory_key="history"
7 )
8
9 conversations = [ # Add conversations
10    ("Tell me about Python", "Python is a versatile language..."),
11    ("What makes it popular?", "Its simplicity and libraries..."),
12    ("Give an example", "Like NumPy for computing..."),
13    ("What about AI?", "Python excels in AI with TensorFlow...")
14 ]
15
16 for input_text, output_text in conversations:
17     memory.save_context({ "input": input_text }, { "output": output_text })
18
19 # Retrieves only last 3 interactions
20 history = memory.load_memory_variables({ })
21 print(f"Stored messages: {len(history['history'])}") # Will be 6 (3 pairs)
```



Document Loaders & Retrievers

YHK

Document Loading

Wide Range of Loaders:

- ▶ **Files:** PDF, Word, PowerPoint, CSV, Markdown
- ▶ **Web:** HTML, URLs, sitemaps, web scraping
- ▶ **Cloud:** S3, GCS, Google Drive, Notion
- ▶ **Databases:** SQL, MongoDB, Elasticsearch
- ▶ **Communication:** Email, Slack, Discord
- ▶ **Code:** GitHub, GitLab, Jupyter notebooks

Modern Document Processing:

```
1 from langchain_community.document_loaders import (
2     PyPDFLoader,
3     WebBaseLoader,
4     TextLoader)
5 from langchain.text_splitters import RecursiveCharacterTextSplitter
6
7 loader = PyPDFLoader("document.pdf")
8 documents = loader.load()
9
10 splitter = RecursiveCharacterTextSplitter(
11     chunk_size=1000,
12     chunk_overlap=200)
13 chunks = splitter.split_documents(documents)
```



Vector Stores & Retrievers

Popular Vector Stores (2024):

- ▶ **Chroma**: Open-source, local-first, easy setup
- ▶ **Pinecone**: Managed service, production-ready
- ▶ **Weaviate**: GraphQL API, hybrid search
- ▶ **Qdrant**: High performance, filtering support
- ▶ **LanceDB**: Serverless, embedded option

Complete Example with Chroma & HuggingFace Embeddings:

```
# from langchain.community.embeddings import HuggingFaceEmbeddings
2 from langchain.huggingface import HuggingFaceEmbeddings
from langchain.community.vectorstores import Chroma
4 from langchain.text_splitters import RecursiveCharacterTextSplitter
from langchain.community.document_loaders import WebBaseLoader
6
loader = WebBaseLoader("https://example.com/article")
8 docs = loader.load()
10 splitter = RecursiveCharacterTextSplitter(chunk_size=1000, chunk_overlap=200)
splits = splitter.split_documents(docs)
12
vectorstore = Chroma.from_documents(
14     documents=splits,
embedding=HuggingFaceEmbeddings(model_name="all-MiniLM-L6-v2"))
16
retriever = vectorstore.as_retriever(search_kwargs={"k": 3})
```



Retrieval Strategies

Different Retrieval Methods:

```
1 # Returns the most semantically similar documents to the user
# query based purely on vector distance (top-k closest matches).
2 retriever = vectorstore.as_retriever(search_type="similarity",
3                                     search_kwargs={"k": 4})
4
5 # MMR (Maximum Marginal Relevance) Balances relevance and diversity,
6 # ensuring the retrieved documents aren't repetitive by reducing semantic redundancy among the top-k results.
7 retriever = vectorstore.as_retriever(search_type="mmr",
8                                     search_kwargs={"k": 4, "fetch_k": 20})
9
10 # Filters out documents that fall below a minimum similarity score,
11 # ensuring only high-confidence matches are returned (still limited by k)
12 retriever = vectorstore.as_retriever(search_type = "similarity_score_threshold",
13                                     search_kwargs={"score_threshold": 0.8, "k": 4})
14
15 # Manual search: Directly retrieves the top-k most similar documents
16 # without using a retriever wrapper same as search_type "similarity" but explicitly called.
17 results = vectorstore.similarity_search("What is machine learning?", k=3)
18
19 # With scores: Same as above but also returns the similarity
20 # score for each document, helping you inspect retrieval quality
21 results_with_scores = vectorstore.similarity_search_with_score(
22     "What is machine learning?", k=3)
```

Chains with LCEL

Modern Chains: LCEL Overview

What Changed?

- ▶ **Old:** LLMChain, SimpleSequentialChain (Deprecated)
- ▶ **New:** LCEL with pipe operator |

LCEL Advantages:

- ▶ **Composability:** Chain components naturally
- ▶ **Streaming:** Built-in streaming support
- ▶ **Async:** Native async/await
- ▶ **Batch:** Process multiple inputs
- ▶ **Fallbacks:** Error handling built-in
- ▶ **Parallelization:** Run steps in parallel
- ▶ **Observability:** Better tracing

Core Runnables:

- ▶ RunnablePassthrough: Pass data through
- ▶ RunnableParallel: Execute in parallel
- ▶ RunnableLambda: Custom functions
- ▶ RunnableBranch: Conditional execution



LCEL: Basic Chain Patterns

```
1 // Simple Chain:  
2 from langchain.groq import ChatGroq  
3 from langchain.core.prompts import ChatPromptTemplate  
4 from langchain.core.output_parsers import StrOutputParser  
5  
6 prompt = ChatPromptTemplate.from_template("Tell me a joke about {topic}")  
7 model = ChatGroq(model.name="gemma-7b-it")  
8 output_parser = StrOutputParser()  
9  
10 chain = prompt | model | output_parser  
11 result = chain.invoke({"topic": "programming"})  
12  
13 // Chain with Multiple Steps:  
14 # Step 1: Generate topic  
15 topic_chain = (  
16     ChatPromptTemplate.from_template("Suggest a {genre} topic")  
17     | ChatGroq(model.name="gemma-7b-it")  
18     | StrOutputParser()  
19 )  
20  
21 # Step 2: Write content  
22 content_chain = (  
23     ChatPromptTemplate.from_template("Write a story about: {topic}")  
24     | ChatGroq(model.name="llama3-8b-8192")  
25     | StrOutputParser()  
26 )  
27  
28 # Combine: topic generates input for content  
29 full_chain = {"topic": topic_chain} | content_chain  
result = full_chain.invoke({"genre": "science fiction"})
```



LCEL: RAG Chain Pattern

Retrieval Augmented Generation:

```
from langchain.core.runnables import RunnablePassthrough
from langchain.groq import ChatGroq
# Assume 'retriever' from previous slide is defined
# Format documents
def format_docs(docs):
    return "\n\n".join(doc.page.content for doc in docs)
# RAG chain
rag_chain = (
    {
        "context": retriever | format_docs,
        "question": RunnablePassthrough()
    }
    | ChatPromptTemplate.from_template("""
        Answer the question based on the context:

        Context: {context}

        Question: {question}
    """)
    | ChatGroq(model_name="llama3-8b-8192")
    | StrOutputParser()
)
# Use it
answer = rag_chain.invoke("What is LangChain?")
```



LCEL: Parallel Execution

Benefits:

- ▶ Faster execution
- ▶ Clean code structure
- ▶ Easy to add/remove tasks

Run Multiple Chains in Parallel:

```
1 from langchain.core.runnables import RunnableParallel
2 from langchain.groq import ChatGroq
3 from langchain.core.prompts import ChatPromptTemplate
4 from langchain.core.output_parsers import StrOutputParser
5
6 # Define parallel tasks
7 parallel_chain = RunnableParallel(
8     summary=ChatPromptTemplate.from_template("Summarize: {text}")
9     | ChatGroq(model_name="llama3-8b-8192")
10    | StrOutputParser(),
11
12    keywords=ChatPromptTemplate.from_template("Extract keywords: {text}")
13    | ChatGroq(model_name="gemma-7b-it")
14    | StrOutputParser(),
15
16    sentiment=ChatPromptTemplate.from_template("Analyze sentiment: {text}")
17    | ChatGroq(model_name="gemma-7b-it")
18    | StrOutputParser()
19 )
20
21 # Execute all in parallel
22 results = parallel_chain.invoke({ "text": "Long document content..." })
```



LCEL: Error Handling & Fallbacks

```
// Fallback to Alternative Model:  
2  from langchain.groq import ChatGroq  
3  from langchain.anthropic import ChatAnthropic  
4  
5  primary = ChatGroq(model.name="llama3-70b-8192")  
6  fallback = ChatAnthropic(model="claude-3-opus-20240229")  
7  
8  chain = prompt | primary.with_fallbacks([fallback]) | output_parser  
9  
10 // Retry Logic: If Groq fails, automatically retries  
11  from langchain.core.runnables import RunnableRetry  
12  
13  chain_with_retry = (  
14      prompt  
15      | RunnableRetry(max_attempts=3, wait_exponential_jitter=True)  
16      | ChatGroq(model.name="gemma-7b-it")  
17      | output_parser  
18  )  
19  
20 // Custom Error Handling:  
21  from langchain.core.runnables import RunnableLambda  
22  
23  def handle_errors(x):  
24      try:  
25          return x  
26      except Exception as e:  
27          return {"error": str(e)}  
28  
chain = prompt | model | RunnableLambda(handle_errors)
```

LCEL: Streaming

```
1 // Stream Tokens as Generated:  
from langchain.groq import ChatGroq  
3 # Assume prompt and output_parser are defined  
chain = prompt | ChatGroq(model_name="gemma-7b-it") | StrOutputParser()  
5  
# Stream output  
7 for chunk in chain.stream({ "topic": "AI" }):  
    print(chunk, end="", flush=True)  
9  
// Async Streaming:  
11 import asyncio  
  
13 async def stream_response():  
    async for chunk in chain.astream({ "topic": "AI" }):  
        print(chunk, end="", flush=True)  
15  
    await asyncio.run(stream_response())  
  
17 // Streaming with Events:  
# Get detailed streaming events  
19 async for event in chain.astream_events({ "topic": "AI" }, version="v1"):  
    kind = event["event"]  
21    if kind == "on_chat_model_stream":  
        content = event["data"]["chunk"].content  
        print(content, end="", flush=True)  
23  
25
```

Agents & Tools

YHK

Modern Agents Overview

What Are Agents?

- ▶ Use LLMs as reasoning engines
- ▶ Dynamically choose which tools to use
- ▶ Iterative: observe, think, act, repeat
- ▶ Best for complex, multi-step tasks

Modern Agent Types (2024):

- ▶ **Tool Calling Agent:** Uses function calling (recommended)
- ▶ **JSON Chat Agent:** Output formatting based on Json
- ▶ **Structured Chat Agent:** For multi-input tools
- ▶ **ReAct Agent:** Reasoning and acting pattern

Deprecated (Don't Use):

- ▶ zero-shot-react-description
- ▶ conversational-react-description
- ▶ self-ask-with-search

(Ref: LangChain Agents Documentation 2024)



Creating Tools: Modern Approach

```
1 // Method 1: Using @tool Decorator:  
2 from langchain.core.tools import tool  
3  
4 @tool  
5 def search_wikipedia(query: str) -> str:  
6     """Search Wikipedia for information about a topic."""  
7     from wikipedia import summary  
8     try:  
9         return summary(query, sentences=3)  
10    except:  
11        return "Could not find information."  
12  
13 @tool  
14 def calculate(expression: str) -> str:  
15     """Evaluate a mathematical expression."""  
16     try:  
17         return str(eval(expression))  
18     except:  
19         return "Invalid expression"  
20  
21 // Method 2: From Function:  
22 \begin{lstlisting}[language=python, basicstyle=\tiny]  
23 from langchain.core.tools import Tool  
24  
25 def get_weather(location: str) -> str:  
26     """Get weather for a location."""  
27     return f"Weather in {location}: Sunny, 72F"  
28  
29 weather_tool = Tool.from_function(  
30     func=get_weather,  
31     name="weather",  
32     description="Get current weather for a location")
```

Tool Calling Agent Example

```
from langchain.groq import ChatGroq
from langchain.agents import create_tool_calling_agent, AgentExecutor
from langchain.core.prompts import ChatPromptTemplate, MessagesPlaceholder
from langchain.core.tools import tool

@tool
def search_wikipedia(query: str) -> str:
    """Search Wikipedia for a topic."""
    # Dummy implementation
    return f"Results for {query} from Wikipedia."

@tool
def calculate(expression: str) -> str:
    """Evaluate a mathematical expression."""
    return str(eval(expression))

tools = [search_wikipedia, calculate]

prompt = ChatPromptTemplate.from_messages([
    ("system", "You are a helpful assistant"),
    ("human", "{input}"),
    MessagesPlaceholder(variable_name="agent_scratchpad")
])

llm = ChatGroq(model.name="llama3-70b-8192", temperature=0)
agent = create_tool_calling_agent(llm, tools, prompt)

agent.executor = AgentExecutor(agent=agent, tools=tools, verbose=True)

result = agent.executor.invoke({"input": "What is the capital of France?"})
```

Modern Approach with bind_tools

```
from langchain.groq import ChatGroq
from langchain.core.tools import tool

@tool
def multiply(a: int, b: int) -> int:
    """Multiply two numbers."""
    return a * b

@tool
def add(a: int, b: int) -> int:
    """Add two numbers."""
    return a + b

# Bind tools to model
llm = ChatGroq(model_name="llama3-8b-8192")
llm.with_tools = llm.bind_tools([multiply, add])

# Invoke
response = llm.with_tools.invoke("What is 3 times 4 plus 5?")

# Check if tool was called
if response.tool_calls:
    tool_call = response.tool_calls[0]
    print(f"Tool: {tool_call['name']}")
    print(f"Args: {tool_call['args']}
```

Output Parsers

YHK

Output Parsers Overview

Why Output Parsers?

- ▶ LLMs return unstructured text
- ▶ Applications need structured data
- ▶ Parsers extract and validate output

Common Parser Types:

- ▶ **StrOutputParser**: Basic string extraction
- ▶ **JsonOutputParser**: Parse JSON responses
- ▶ **PydanticOutputParser**: Structured data with validation
- ▶ **StructuredOutputParser**: Multiple fields
- ▶ **CommaSeparatedListOutputParser**: Lists

Modern Best Practice:

- ▶ Use OpenAI's `with_structured_output()` when possible
- ▶ More reliable than prompt-based parsing
- ▶ Leverages function calling internally

Structured Output: Modern Approach

Benefits:

- ▶ Type-safe responses
- ▶ Automatic validation
- ▶ More reliable than prompt-based parsing

```
1 from langchain.groq import ChatGroq
  from pydantic import BaseModel, Field
3
# Define output schema
5 class Person(BaseModel):
    name: str = Field(description="Person's name")
7    age: int = Field(description="Person's age")
  occupation: str = Field(description="Person's job")
9
# Create model with structured output (relies on tool-calling)
11 llm = ChatGroq(model.name="llama3-70b-8192")
  structured_llm = llm.with_structured_output(Person)
13
# Use in chain
15 response = structured_llm.invoke(
    "Tell me about a software engineer named Alice who is 28"
)
17
# Response is a Pydantic object
19 print(response.name)
21 print(response.age)
  print(response.occupation)
```



Pydantic Output Parser (Alternative)

```
from langchain.output_parsers import PydanticOutputParser
from langchain.groq import ChatGroq
from pydantic import BaseModel, Field, validator
from typing import List

class MovieReview(BaseModel):
    title: str = Field(description="Movie title")
    rating: int = Field(description="Rating from 1–10")

parser = PydanticOutputParser(pydantic.object=MovieReview)

# Add to prompt
prompt = ChatPromptTemplate.from_template("""
Review the movie: {movie}

{format_instructions}""")

chain = (prompt.partial(format_instructions=parser.get_format_instructions()
    | ChatGroq(model.name="gemma-7b-it") | parser)
result = chain.invoke({ "movie": "Inception" })
```

Output Parser Error Handling

```
1 from langchain.groq import ChatGroq
# Assume 'parser' is a defined PydanticOutputParser
3
// OutputFixingParser (Auto-fix Errors):
5 from langchain.output_parsers import OutputFixingParser
7
# If parsing fails, use LLM to fix
fixing_parser = OutputFixingParser.from_llm(
9     parser=parser,
    llm=ChatGroq(model.name="gemma-7b-it")
11 )
13
// Automatically fixes malformed output
# result = fixing_parser.parse(malformed_output)
15
// RetryOutputParser (Retry with Context):
17 from langchain.output_parsers import RetryWithErrorOutputParser
19
retry_parser = RetryWithErrorOutputParser.from_llm(
    parser=parser,
    llm=ChatGroq(model.name="llama3-8b-8192")
21 )
23
# Retries with both output and original prompt
25 # result = retry_parser.parse_with_prompt(...)
```



LangChain Ecosystem

LangChain Ecosystem Components

- ▶ **LangChain Core:**
 - ▶ Base abstractions and LCEL
 - ▶ Foundation for all other packages
- ▶ **LangChain Community:**
 - ▶ Third-party integrations
 - ▶ Vector stores, document loaders
 - ▶ Community-maintained tools
- ▶ **LangGraph:**
 - ▶ Build stateful, multi-actor applications
 - ▶ Complex agent workflows with cycles
 - ▶ State management for agents
- ▶ **LangServe:**
 - ▶ Deploy chains as REST APIs
 - ▶ Automatic FastAPI generation
 - ▶ Production deployment
- ▶ **LangSmith:**
 - ▶ Debugging and monitoring
 - ▶ Tracing and evaluation
 - ▶ Dataset management

LangGraph: Stateful Agents

- ▶ Build complex, stateful agent workflows
- ▶ Support for cycles and conditional logic
- ▶ Persist state across interactions
- ▶ Multiple agents working together

```
1 from langgraph.graph import StateGraph
2 from typing import TypedDict, Annotated
3 import operator
4
5 class AgentState(TypedDict):
6     messages: Annotated[list, operator.add]
7     next: str
8
9     def call_model(state):
10         response = llm.invoke(state["messages"])
11         return { "messages": [response] }
12
13     def should_continue(state):
14         if len(state["messages"]) > 5:
15             return "end"
16         return "continue"
17
18 workflow = StateGraph(AgentState)
19 workflow.add_node("agent", call_model)
20 workflow.add_conditional_edges("agent", should_continue)
21 workflow.set_entry_point("agent")
22
23 app = workflow.compile()
```

LangServe: Deploy as API

Automatic Features:

- ▶ Interactive playground at /chat/playground
- ▶ OpenAPI docs at /docs
- ▶ Streaming support
- ▶ Batch processing

Deploy Any Chain as REST API:

```
1 from fastapi import FastAPI
2 from langchain import add_routes
3 from langchain_groq import ChatGroq
4 from langchain_core.prompts import ChatPromptTemplate
5
6 # Create chain
7 prompt = ChatPromptTemplate.from_template("Tell me about {topic}")
8 chain = prompt | ChatGroq(model_name="gemma-7b-it")
9
10 # Create FastAPI app
11 app = FastAPI(
12     title="LangChain API", version="1.0",
13     description="API for my LangChain application")
14
15 # Add chain as route
16 add_routes(app, chain, path="/chat")
17
18 # Run: uvicorn app:app --reload
19 # API available at: http://localhost:8000/chat
```



LangSmith: Monitoring & Debugging

What is LangSmith?

- ▶ Platform for debugging LLM applications
- ▶ Trace every step of chain execution
- ▶ Evaluate model performance
- ▶ Manage test datasets
- ▶ Monitor production applications

Setup:

```
1 import os
2 from langchain.groq import ChatGroq
3 from langchain.core.prompts import ChatPromptTemplate
4 from langchain.core.output_parsers import StrOutputParser
5
6 os.environ["LANGCHAIN_TRACING_V2"] = "true"
7 os.environ["LANGCHAIN_API_KEY"] = "your-langsmith-api-key"
8 os.environ["LANGCHAIN_PROJECT"] = "my-groq-project"
9
10 # Now all chain executions are automatically traced
11 prompt = ChatPromptTemplate.from_template("Tell me about {topic}")
12 llm = ChatGroq(model.name="gemma-7b-it")
13 output_parser = StrOutputParser()
14 chain = prompt | llm | output_parser
15 result = chain.invoke({"topic": "Large Language Models"})
16
17 # View traces at: https://smith.langchain.com
```



Best Practices

YHK

Best Practices: Error Handling

```
1 // 1. Use Fallbacks:  
2 from langchain.groq import ChatGroq  
3 from langchain.anthropic import ChatAnthropic  
4  
5 primary = ChatGroq(model.name="llama3-70b-8192")  
6 fallback = ChatAnthropic(model="claude-3-opus-20240229")  
7  
8 chain = prompt | primary.with_fallbacks([fallback]) | parser  
9  
// 2. Implement Retries:  
10 from langchain.core.runnables import RunnableRetry  
11  
12 chain_with_retry = (  
13     prompt  
14     | RunnableRetry(  
15         max_attempts=3,  
16         wait_exponential_jitter=True  
17     )  
18     | ChatGroq(model.name="gemma-7b-it")  
19     | parser  
20 )  
21  
22 // 3. Graceful Degradation:  
23 try:  
24     result = chain.invoke(input_data)  
25 except Exception as e:  
26     # logger.error(f"Chain failed: {e}")  
27     # result = fallback.response()  
28     pass
```

Best Practices: Token Management with Groq

Groq Models - Fast and Free Tier:

- ▶ gemma2-9b-it: Ultra-fast, 8K context, great for chat
- ▶ llama3-8b-8192: Balanced, 8K context, general purpose
- ▶ llama3-70b-8192: Most capable, 8K context, complex tasks
- ▶ llama-3.1-8b-instant: Fastest, optimized for low latency
- ▶ Groq billed by requests/day on free tier, not tokens

```
1 from langchain.groq import ChatGroq
2
3 # Choose model based on needs
4 llm_fast = ChatGroq(
5     model_name="gemma2-9b-it",
6     max_tokens=500,
7     temperature=0.7,
8     max_retries=2
9 )
10
11 # Monitor context length manually
12 from transformers import AutoTokenizer
13 tokenizer = AutoTokenizer.from_pretrained("google/gemma-2-9b-it")
14 token_count = len(tokenizer.encode(text))
15
16 # Groq's speed allows batch processing without latency concerns
```



Best Practices: Security & Privacy

- ▶ Be aware of the data usage policies of your LLM provider.
- ▶ Groq has a zero-retention policy for API data.
- ▶ Consider self-hosted models for maximum data control.
- ▶ Implement PII detection and redaction before sending data.

```
// 1. API Key Management:  
1 import os  
2 from dotenv import load_dotenv  
3  
4 load_dotenv() # Load from .env file  
5 api_key = os.getenv("GROQ_API_KEY")  
6  
// 2. Timeouts and Retries:  
7 from langchain.groq import ChatGroq  
8  
10 llm = ChatGroq(model.name="llama3-8b-8192",  
11     temperature=0,  
12     max_retries=2,  
13     request_timeout=30) # seconds  
14  
// 3. Input Validation:  
15 def validate_input(user_input: str) -> str:  
16     if len(user_input) > 8000: # Sanitize input  
17         raise ValueError("Input too long for the model context.")  
18     # Remove potential injection attempts  
19     return user_input.strip()  
20
```



Best Practices: Choosing the Right Model

Groq Model Selection Guide:

Use Case	Model	Why
Chatbots	gemma2-9b-it	Fast, efficient, good reasoning
Summarization	llama3-8b-8192	Balanced speed/quality
Complex reasoning	llama3-70b-8192	Most capable
Ultra-low latency	llama-3.1-8b-instant	Optimized for speed
Code generation	llama3-70b-8192	Better instruction following

```
1 # Pattern: Start with fast model, fallback to capable model
2 from langchain.groq import ChatGroq
3
4 primary = ChatGroq(model_name="gemma2-9b-it", temperature=0.7)
5 fallback = ChatGroq(model_name="llama3-70b-8192", temperature=0.7)
6
7 chain = prompt | primary.with_fallbacks([fallback]) | parser
```

Best Practices: Async Patterns

Benefits:

- ▶ Faster parallel processing
- ▶ Better resource utilization
- ▶ Improved user experience with streaming

```
1 // Use Async for Better Performance:  
2 import asyncio  
3  
4 async def process_multiple_queries(queries):  
5     # Process queries concurrently  
6     tasks = [chain.ainvoke({ "input": q }) for q in queries]  
7     results = await asyncio.gather(*tasks)  
8     return results  
9  
10 # Run  
11 queries = ["Query 1", "Query 2", "Query 3"]  
12 results = asyncio.run(process_multiple_queries())  
13 // Async Streaming:  
14 async def stream_response(input_text):  
15     async for chunk in chain.astream( { "input": input_text } ):  
16         print(chunk, end="", flush=True)  
17  
18 asyncio.run(stream_response("Tell me about AI"))
```

What's New in LangChain Ecosystem

(Oct 2025, Relase of 1.0 version)



Why LangGraph 1.0 is Nearly Non-Breaking and Production-Ready

- ▶ LangGraph 1.0 introduces minimal breaking changes from prior releases.
- ▶ Existing LangGraph implementations continue to work as-is.
- ▶ Core runtime stable across Python and TypeScript.
- ▶ Supports durable execution with checkpoints and rollback.
- ▶ State can persist via Postgres or SQLite for reliability.
- ▶ Human-in-the-loop and interrupt handling built-in.
- ▶ Ready for production-grade use; used by Uber, LinkedIn, Klarna, JPMorgan, Cloudflare.



How LangChain 1.0 Simplifies Agents to Just 10 Lines of Code

- ▶ LangChain 1.0 Alpha released; full 1.0 expected by late October.
- ▶ Rewritten as a simplified agent runtime built on LangGraph.
- ▶ Agents can now be created in roughly 10 lines of code.
- ▶ Legacy LangChain moved to “LangChain Classic.”
- ▶ Functionally similar to OpenAI or Pedantic SDK agents.
- ▶ Unified “create_agent” abstraction across languages.
- ▶ Integrates easily with external SDKs like Google’s 80k.

Durable Execution, Streaming, Human-in-the-Loop, and Time Travel

- ▶ Agents persist state with rollback and checkpointing.
- ▶ Supports human approval and manual intervention mid-run.
- ▶ Time travel enables returning to earlier workflow states.
- ▶ Four streaming modes: messages, updates, values, custom.
- ▶ Update streaming ideal for dashboards and UX refresh.
- ▶ State persistence possible locally or via Postgres.
- ▶ Enables retry and branch execution for debugging and auditing.

The New Standardized Content Blocks for Easier Model Switching

- ▶ LangChain Core now uses standardized content/message blocks.
- ▶ Abstracts input/output across providers like OpenAI and Anthropic.
- ▶ Simplifies switching models without rewriting logic.
- ▶ Unifies multimodal inputs, reasoning traces, and tool calls.
- ▶ Middleware layer ensures consistent formatting and metadata.
- ▶ Supports normalized logging across different model providers.
- ▶ Key enabler for multi-model experimentation and portability.



Real-World Use Cases

- ▶ Meeting notes enrichment and summarization pipelines.
- ▶ Financial EDI approval flows with human-in-loop steps.
- ▶ Real-time dashboards using streaming “update” mode.
- ▶ Persistent agent state across cloud or local environments.
- ▶ Dynamic model selection for efficiency (e.g., switching from Claude to Llama3).
- ▶ Internal RAG/Graph-RAG systems leveraging PGVector.
- ▶ Used at scale in production by enterprises and startups alike.



Example: Defining a Simple Agent in LangChain 1.0

```
from langchain.agents import create_agent
2
agent = create_agent(
    tools=[search, summarize],
    model="gpt-4-turbo"
)
6
8 result = agent("Summarize today's meeting notes")
print(result)
```

Anti-Patterns to Avoid

- ▶ Token-only streaming without structured updates.
- ▶ Relying on ephemeral in-memory state, persist externally.
- ▶ Single-function “flat” agents instead of structured graphs.
- ▶ Provider lock-in design for model and SDK portability.
- ▶ Neglecting checkpointing and recovery strategies.
- ▶ Ignoring observability, use LangSmith, LangFuse, or OpenTelemetry.
- ▶ Skipping standardized content blocks creates fragility.

Conclusions

YHK

What is LangChain?

LangChain: A Comprehensive Framework for Building LLM-Powered Applications

- ▶ **Core Purpose:** Simplify development of complex applications using Large Language Models
- ▶ **Key Components:**
 - ▶ Unified interface for multiple language models
 - ▶ Advanced prompt management
 - ▶ Flexible data integration
 - ▶ Intelligent agent and tool systems
- ▶ **Founding:** Created by Harrison Chase as an open-source project
- ▶ **Availability:** Python and JavaScript libraries

Feature	Description
Data Awareness	Connect LLMs to external data sources
Agentic Capability	Enable dynamic interaction with environment
Model Flexibility	Support for multiple LLM providers

(Ref: LangChain Framework Overview)

Why LangChain?

Addressing Challenges in LLM Application Development

► Limitations in Current LLM Tooling:

- Fragmented model interfaces
- Complex prompt management
- Limited external data integration
- Lack of flexible reasoning mechanisms

► LangChain Solutions:

- Model-agnostic framework
- Standardized prompt engineering
- Seamless external data augmentation
- Intelligent agent orchestration

► Key Advantages:

- Rapid prototyping of AI applications
- Simplified model and tool integration
- Scalable architecture for complex use cases

Enabling Developers to Build Sophisticated AI Systems with Ease

(Ref: LangChain Development Principles)



One pager

- ▶ **Models:** Building blocks supporting different AI model types - LLMs, Chat, Text Embeddings.
- ▶ **Prompts:** Inputs constructed from various components. LangChain offers easy interfaces - Prompt Templates, Example Selectors, Output Parsers.
- ▶ **Memory:** Stores/retrieves messages, short or long term, in conversations.
- ▶ **Indexes:** Assist LLMs with documents - Document Loaders, Text Splitters, Vector Stores, Retrievers.
- ▶ **Chains:** Combine components or chains in order to accomplish tasks.
- ▶ **Agents:** Empower LLMs to interact with external systems, make decisions, and complete tasks using Tools.



(Ref: Building Generative AI applications made easy with Vertex AI PaLM API and LangChain - Anand Iyer, Rajesh Thallam)

Resources & Next Steps

Official Documentation:

- ▶ LangChain Docs: <https://python.langchain.com>
- ▶ LangChain Blog: <https://blog.langchain.dev>
- ▶ LangChain Academy: <https://academy.langchain.com>
- ▶ API Reference: <https://api.python.langchain.com>

GitHub Repositories:

- ▶ Core: <https://github.com/langchain-ai/langchain>
- ▶ Templates:
<https://github.com/langchain-ai/langchain/tree/master/templates>
- ▶ LangGraph: <https://github.com/langchain-ai/langgraph>

Community:

- ▶ Discord: <https://discord.gg/langchain>
- ▶ Twitter: @LangChainAI
- ▶ YouTube: LangChain official channel

Practice:

- ▶ Start with simple LCEL chains
- ▶ Build a RAG application
- ▶ Create custom tools and agents
- ▶ Deploy with LangServe

Closing Thoughts

- ▶ Langchain's usefulness in solving problems today
- ▶ Possibility of LLM APIs expanding capabilities over time
- ▶ Potential for Langchain to become an interface to LLMs
- ▶ Acknowledgment of Langchain's valuable contributions and community efforts
- ▶ Appreciation for the work done by Harrison and the Langchain team

(Ref: What is Langchain and why should I care as a developer? - Logan Kilpatrick)

References

Many publicly available resources have been refereed for making this presentation. Some of the notable ones are:

- ▶ Intro to LangChain for Beginners - A code-based walkthrough- Menlo Park Lab
- ▶ LangChain Crash Course (10 minutes): Easy-to-Follow Walkthrough of the Most Important Concepts - Menlo Park Lab
- ▶ Official doc: <https://docs.langchain.com/docs/>
- ▶ Git Repo: <https://github.com/hwchase17/langchain>
- ▶ LangChain 101: The Complete Beginner's Guide Edrick <https://www.youtube.com/watch?v=P3MAbZ2eMUI> A wonderful overview, don't miss
- ▶ Cookbook by Gregory Kamradt(Easy way to get started): <https://github.com/gkamradt/langchain-tutorials/blob/main/LangChain%20Cookbook.ipynb>
- ▶ Youtube Tutorials: https://www.youtube.com/watch?v=_v_fgW2SkkQ
- ▶ LangChain 101 Course (updated 2024 with LCEL) - Ivan Reznikov
- ▶ <https://github.com/IvanReznikov/DataVerse/tree/main/Courses/LangChain>
- ▶ A Complete LangChain Guide <https://nanonets.com/blog/langchain/>
- ▶ 7 Ways to Use LangSmith's Superpowers to Turboboost Your LLM Apps - Menlo Park Lab
- ▶ LangChain official blog: <https://blog.langchain.dev>
- ▶ LangChain Academy: <https://academy.langchain.com>
- ▶ LangChain templates repo: <https://github.com/langchain-ai/langchain/tree/master/templates>
- ▶ Groq Documentation: <https://console.groq.com/docs>
- ▶ Groq Model Playground: <https://groq.com/>

Thanks ...

- ▶ Office Hours: Saturdays, 3 to 5 pm (IST);
Free-Open to all; email for appointment to
yogeshkulkarni at yahoo dot com
- ▶ Call + 9 1 9 8 9 0 2 5 1 4 0 6



(<https://www.linkedin.com/in/yogeshkulkarni/>)



(<https://medium.com/@yogeshharibhaukul>)



(<https://www.github.com/yogeshhk/>)

Pune AI Community (PAIC)

- ▶ Two-way communication:
 - ▶ Website puneaicommunity dot org
 - ▶ Email puneaicommunity at gmail dot com
 - ▶ Call + 9 1 9 8 9 0 2 5 1 4 0 6
 - ▶ LinkedIn:
<https://linkedin.com/company/pune-ai-community>
- ▶ One-way Announcements:
 - ▶ Twitter (X) @puneaicommunity
 - ▶ Instagram @puneaicommunity
 - ▶ WhatsApp Community: Invitation Link
<https://chat.whatsapp.com/LluOrhyEzuQLDr25ixZ>
 - ▶ Luma Event Calendar: puneaicommunity
- ▶ Contribution Channels:
 - ▶ GitHub: Pune-AI-Community and puneaicommunity
 - ▶ Medium: pune-ai-community
 - ▶ YouTube: @puneaicommunity



Website

Pune AI Community (PAIC) QR codes



Website



Medium Blogs



Twitter-X



LinkedIn Page



Github Repository



WhatsApp Invite



Luma Events



YouTube Videos



Instagram

