

MULTI-AGENT DECISION SYSTEMS

Yogesh Haribhau Kulkarni



Outline

① INTRO

② MULTI-AGENT

③ REFS

About Me

YHK

Yogesh Haribhau Kulkarni

Bio:

- ▶ 20+ years in CAD/Engineering software development
- ▶ Got Bachelors, Masters and Doctoral degrees in Mechanical Engineering (specialization: Geometric Modeling Algorithms).
- ▶ Currently doing Coaching in fields such as Data Science, Artificial Intelligence Machine-Deep Learning (ML/DL) and Natural Language Processing (NLP).
- ▶ Feel free to follow me at:
 - ▶ Github (github.com/yogeshhk)
 - ▶ LinkedIn (www.linkedin.com/in/yogeshkulkarni/)
 - ▶ Medium (yogeshharibhaukulkarni.medium.com)
 - ▶ Send email to [yogeshkulkarni at yahoo dot com](mailto:yogeshkulkarni@yahoo.com)



Office Hours:
Saturdays, 2 to 5pm
(IST); Free-Open to all;
email for appointment.

Introduction

YHK

Classical idea of Agents

- ▶ Agents are autonomous
- ▶ Can look at their environment and analyze the situation
- ▶ Make comprehensive plans to achieve specific goals
- ▶ Actually take action to execute those plans
- ▶ Agents bridge the gap between answering and doing

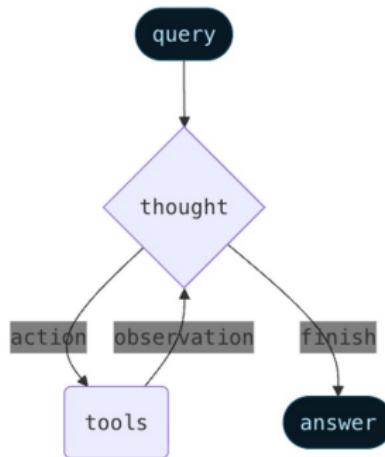
Welcome to AI Agents

- ▶ Agents were there from 1950's but they are effective because of LLMs.
- ▶ Agents are systems where LLMs dynamically direct their own paths and tool usage
- ▶ Agents can have autonomous-ness or predefined workflow paths
- ▶ Essential component in modern AI systems with varying degrees of autonomy
- ▶ Unlike LLMs that just respond to prompts, agents do things
- ▶ Not just everyday chatbots, systems that reason, plan, and take action
- ▶ Can take on complex multi-step tasks autonomously or with human-in-loop
- ▶ Technology is advancing rapidly from conversational to agentic AI
- ▶ AI agents represent one of the most exciting frontiers in AI
- ▶ 2025 is the year of AI agents.

The Impossible Job

- ▶ If your boss tasks you with planning a massive get-together
- ▶ Must prepare a guest list and plan fancy menu
- ▶ Needs to find entertainment for the event
- ▶ A simple chatbot (with answering LLM, not reasoning LLM) cannot handle these complex requirements
- ▶ You need an AI agent, not just responses, but actions
- ▶ Agents have a Reasoning LLM as brain and tools as its hands-legs.
- ▶ AI agents improve with reasoning, acting, and memory components.
(ReAct = Reasoning + Acting)

ReAct (Reasoning + Acting) Agent



A ReAct agent selects tools by thinking about the request, selecting a tool, observing its result, and iterating until the request is fulfilled.

Based on paper:

[ReAct: Synergizing Reasoning and Acting in Language Models](#)

<https://arxiv.org/abs/2210.03629>

Diagram from Langchain docs: <https://docs.langchain.com/oss/python/langchain-agents>

(Ref: Python + AI Agents - Microsoft)

The Evolution of AI Capabilities

- ▶ **Traditional Programming:** Needed code to operate
- ▶ **Traditional ML:** Needed feature engineering
- ▶ **Deep Learning:** ML with Neural networks, no feature engineering
- ▶ **Foundational Models:** Many tasks single model
- ▶ **Agents (2024 . . .):** Can actually **do things**, not just talk

Why Does “Taking Action” Matter?

- ▶ In 2022, ChatGPT was revolutionary because AI felt conversational
- ▶ By 2024, people wanted more than conversation, they wanted **execution**
- ▶ Examples of what users now expect:
 - ▶ Instead of listing leads ? **email them directly**
 - ▶ Instead of summarizing docs ? **file and create workflow tasks**
 - ▶ Instead of suggesting products ? **customize landing pages**
- ▶ This shift from **information** to **action** defines the agent era

How Agents Work?

- ▶ Agent acts, take you from one state to the other state(ReAct paper: Reasoning and Action),
- ▶ It can plan and make decisions, provides value by workflow automation.
- ▶ Agents have access to tools (ToolFormer paper) e.g. Search APIs, booking, send email etc.
- ▶ Interacting of external environment and other Agents, etc.
- ▶ Memory to keep the history of conversations/actions done so far.
- ▶ May have human-in-loop to keep it sane in the wild-world.

The Agent's Fundamental Game Loop

- ▶ Not a one-and-done action, but a continuous reasoning loop
- ▶ Similar to a programming while loop that keeps iterating
- ▶ **Thought:** Analyzes situation and plans next step
- ▶ **Action:** Calls specific tools to execute the plan
- ▶ **Observation:** Examines results of the action taken
- ▶ Cycle repeats: Thought → Action → Observation
- ▶ Continues until the task is completely accomplished
- ▶ This loop enables continuous adaptation and problem-solving

Agent's Inner Monologue

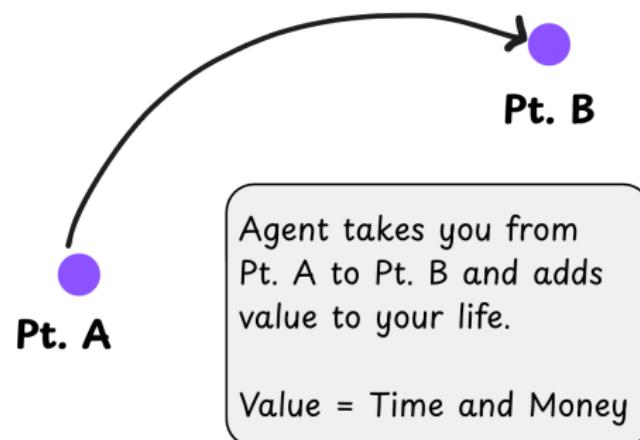
- ▶ Agents have visible thought processes before taking action
- ▶ Example: "User wants weather in New York. I have a tool for that".
- ▶ "My first move is to call the weather API".
- ▶ Internal planning step makes agents more than reactive programs
- ▶ Reasoning through problems before execution
- ▶ This deliberation distinguishes agents from simple scripts
- ▶ Shows intelligent decision-making rather than blind execution

How Do Agents Take Action?

- ▶ The magic lies in **tools** and **function calling**
- ▶ Agents are paired with APIs, plugins, or external systems
- ▶ Instead of just text responses, LLMs output structured commands:
 - ▶ “Call the send_email() function with these inputs...”
 - ▶ “Fetch records from CRM using this query...”
 - ▶ “Schedule a meeting for Tuesday at 2PM...”
- ▶ **Mental model:** LLM = brain, Tools = hands
- ▶ Without tools, agents just talk. With tools, they act.

Defining AI Agents with an example

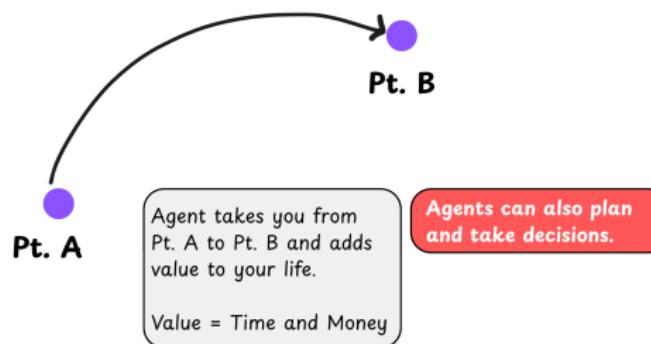
- ▶ Planning a trip involves many complex tasks
- ▶ Point A: Just discussing the trip
- ▶ Point B: All bookings and itinerary ready
- ▶ AI Agents aim to take you from A to B
- ▶ First requirement: Agent adds value by saving time/money



(Ref: Vizuara AI Agents Bootcamp)

Evolving Definition of Agents

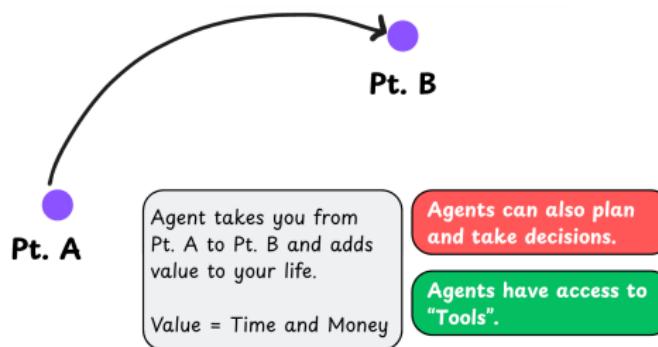
- ▶ Not all tools from A to B are agents (e.g., cars)
- ▶ Agents must plan and make decisions
- ▶ Example: Choosing flights based on budget
- ▶ Planning daily itinerary needs contextual judgment
- ▶ Second requirement: Agent includes decision-making ability



(Ref: Vizuara AI Agents Bootcamp)

Agents Need Tools

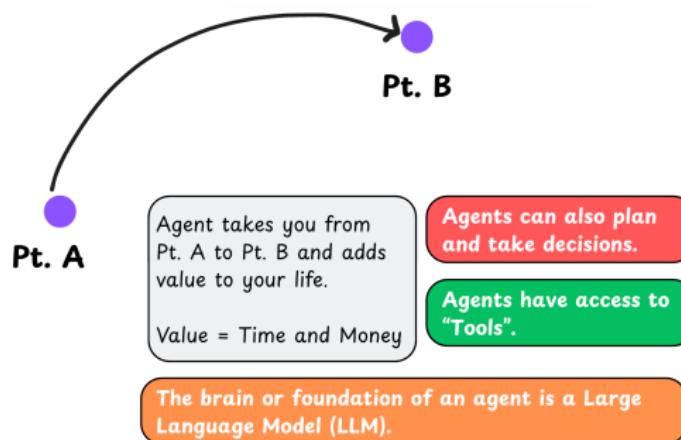
- ▶ Even self-driving cars plan but are not agents
- ▶ Agents need access to external tools
- ▶ Tools = Access to services (e.g., Gmail, Booking)
- ▶ Agents perform tasks using these tools
- ▶ Third requirement: Agent adds tool access to capabilities



(Ref: Vizuara AI Agents Bootcamp)

Rise of LLMs in Agents

- ▶ Transformers (2017) enabled powerful LLMs
- ▶ LLMs understand and generate human language
- ▶ Agents use LLMs for reasoning and planning
- ▶ LLMs enable understanding of webpages and writing emails
- ▶ Fourth requirement: Agents are LLMs with tools and planning ability



(Ref: Vizuara AI Agents Bootcamp)

What Is an Agent? (Technical Definition)

- ▶ Agent acts and takes you from one state to another, providing value through workflow automation
- ▶ Based on ReAct paradigm: **Reasoning + Acting**
- ▶ Key capabilities:
 - ▶ Can plan and make decisions
 - ▶ Has access to tools (search APIs, booking, email, etc.)
 - ▶ Interacts with external environments and other agents
 - ▶ Maintains memory of conversations and actions
 - ▶ May include human-in-the-loop for safety
- ▶ Agents existed since the 1950s but are now effective because of LLMs

Two Ways to Define Agents

Technical View:

- ▶ LLM (brain)
- ▶ + Tools (hands)
- ▶ + Planning (strategy)
- ▶ + Memory (context)
- ▶ + State management

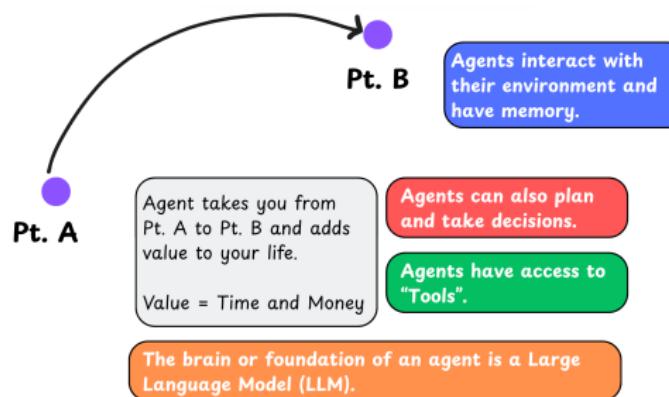
Business View:

- ▶ Systems that complete tasks end-to-end
- ▶ Focus on outcomes, not components
- ▶ Solve real-world problems
- ▶ Provide measurable value

Important: Today's agents are **engineering wrappers** around AI models, the intelligence comes from the LLMs, agents help act on that intelligence.

Final Definition of Agents

- ▶ Agents can learn from feedback and environment
- ▶ Agents interact with tools, humans, and websites
- ▶ They improve with experience (memory)
- ▶ Agents evolve over time via memory and feedback
- ▶ Fifth requirement: LLMs + Tools + Planning + Learning



(Ref: Vizuara AI Agents Bootcamp)

Understanding Agency

- ▶ Agency = Level of autonomy an agent has
- ▶ Low agency → less value
- ▶ High agency → high value
- ▶ More autonomous agents can handle complex tasks
- ▶ Agency is key to measuring agent usefulness

| Agency Level | Description | Name | Example | 🔗 |
|--------------|--|------------------|---|---|
| ●○○○○ | Agent does not influence what happens next | Simple Processor | Grammar checker that rewrites sentences | |
| ●●○○○ | Agent determines basic control flow | Router | Customer Query → Tech Support or Sales | |
| ●●●○○ | Agent determines function execution | Tool Caller | A smart calendar assistant that spots "let's meet on Tuesday" and books the meeting | |
| ●●●●○ | Lays out a short plan and carries it step by step | Multi-step Agent | Personal travel planner that gathers flight options, hotels, local activities | |
| ●●●●● | One agentic workflow starts another agentic workflow | Multi-agent | Travel planner agent → Booking agent ← Email agent | |

(Ref: Vizuara AI Agents Bootcamp)



Tools: The Agent's Hands

- ▶ LLM is the agent's brain; tools are its hands
- ▶ Tools are functions agents call to interact with the world
- ▶ Can search web, run calculations, or query databases
- ▶ Bridge between thinking and doing in the real world
- ▶ Enable agents to move from planning to execution
- ▶ Without tools, agent thoughts would be useless
- ▶ Tools provide the interface to external systems and data

Two Ways Agents Use Tools

- ▶ **JSON Agent:** Writes structured work orders for other systems
- ▶ JSON approach requires external system to read and execute
- ▶ **Code Agent:** Directly writes and runs code blocks
- ▶ Code approach is more direct and powerful
- ▶ Code is naturally more expressive than JSON
- ▶ Can handle complex logic like loops and conditionals
- ▶ Modular, easier to debug, and taps into existing libraries
- ▶ Code agents can access thousands of APIs directly

Code Agent in Action

- ▶ Alfred needs a gala menu - agent has “suggest_menu” tool
- ▶ Agent doesn't just make up suggestions randomly
- ▶ Generates and runs actual code to call the specific tool
- ▶ Gets real results from the tool execution
- ▶ Super direct, efficient, and powerful way to take action
- ▶ Code generation enables precise tool interaction
- ▶ Results are based on actual tool capabilities, not hallucination



Advanced Pattern: Agentic RAG

- ▶ Traditional RAG: Retrieval Augmented Generation fetches info before answering
- ▶ Agentic RAG supercharges this with intelligent multi-step processes
- ▶ Turns retrieval itself into an agent-driven task
- ▶ Like having a master researcher on staff
- ▶ Doesn't just do one search - runs complete research processes
- ▶ Rewrites queries for better results and runs multiple searches
- ▶ Uses findings to inform next searches and validates accuracy
- ▶ Pulls from both private data and public web sources

Multi-Agent Systems: Digital Teams

- ▶ Complex problems like finding the missing Batmobile need teams
- ▶ Single agents can't handle web searches, calculations, and visualization
- ▶ Solution: Build teams of specialized agents
- ▶ Manager agent acts as project lead breaking down big tasks
- ▶ Delegates work to specialist agents with specific skills
- ▶ Web agent handles online searching while manager coordinates
- ▶ Manager focuses on big picture and final integration
- ▶ Digital division of labor for complex problem solving

Papers that Shaped AI Agents

- ▶ Core research papers laid the foundation
- ▶ Introduced key frameworks and architectures
- ▶ Sparked recent boom in agent development
- ▶ Include Transformer and Agentic frameworks
- ▶ Major driving force in LLM-based agent systems

Chain-of-Thought Prompting Elicits Reasoning in Large Language Models

Jason Wei Xuezhi Wang Dale Schuurmans Maarten Bosma

Brian Ichter Fei Xia Ed H. Chi Quoc V. Le Denny Zhou

Google Research, Brain Team
{jasonwei,dennyzhou}@google.com

REACT: SYNERGIZING REASONING AND ACTING IN LANGUAGE MODELS

Shunyu Yan¹, Jeffrey Zhao², Dian Yu², Nan Du², Izhak Shafran², Karthik Narasimhan¹, Yuan Cao²

¹Department of Computer Science, Princeton University

²Google Research, Brain team

¹{shunuyu,karthikn}@princeton.edu

²{jeffreyzhao,diyanyu,dunan,izhak,yuancao}@google.com

Toolformer: Language Models Can Teach Themselves to Use Tools

Timo Schick Jane Dwivedi-Yu Roberto Dessì¹ Roberta Raileanu
Maria Lomeli Luke Zettlemoyer Nicola Cancedda Thomas Scialom

Meta AI Research ¹Universitat Pompeu Fabra

Generative Agents: Interactive Simulacra of Human Behavior

Joon Sung Park
Stanford University
Stanford, USA
joonspk@stanford.edu

Joseph C. O'Brien
Stanford University
Stanford, USA
jobjen3@stanford.edu

Carrie J. Cai
Google Research
Mountain View, CA, USA
cja@ai.google.com

Meredith Ringel Morris
Google DeepMind
Seattle, WA, USA
mmemo@google.com

Percy Liang
Stanford University
Stanford, USA
pliang@cs.stanford.edu

Michael S. Bernstein
Stanford University
Stanford, USA
mbs@cs.stanford.edu

(Ref: Vizuara AI Agents Bootcamp)



Frameworks

YHK

Python AI Agent Frameworks Overview

- ▶ Python offers several frameworks for building single and multi-agent AI systems.
- ▶ These frameworks support workflows, tool use, reasoning, and human-in-the-loop capabilities.
- ▶ Common applications include autonomous assistants, research copilots, and workflow orchestration.



LangChain v1

- ▶ **Type:** Multi-agent, graph-based framework.
- ▶ **Built on:** LangGraph for agent state management and orchestration.
- ▶ **Features:**
 - ▶ Agent-centric architecture.
 - ▶ Workflow graphs for reasoning and tool execution.
 - ▶ Human-in-the-loop integration via LangSmith.
 - ▶ Supports OpenAI, Anthropic, Gemini, and local models.

Pydantic-AI

- ▶ **Type:** Lightweight single/multi-agent framework.
- ▶ **Focus:** Strong type safety using Pydantic for structured inputs/outputs.
- ▶ **Features:**
 - ▶ Declarative agent definition with validation.
 - ▶ Easy LLM orchestration and chaining.
 - ▶ Built for developers needing data integrity.
- ▶ **Ideal for:** AI applications requiring schema consistency and typed reasoning.



Crew AI

- ▶ **Type:** Multi-agent coordination framework.
- ▶ **Concept:** “Crew” of specialized agents working collaboratively.
- ▶ **Features:**
 - ▶ Role-based agent interactions.
 - ▶ Shared memory and task assignment.
 - ▶ Supports autonomous workflow creation.
 - ▶ Integrates with LangChain, OpenAI, and custom tools.
- ▶ **Use Case:** Complex multi-role task solving (research, planning, content generation).

Autogen

- ▶ **Developed by:** Microsoft Research.
- ▶ **Type:** Multi-agent conversational framework.
- ▶ **Features:**
 - ▶ Agent-to-agent and human-in-loop communication.
 - ▶ Supports function calling and code execution.
 - ▶ Built-in memory and orchestration logic.
 - ▶ Extendable with custom agents and tools.
- ▶ **Example Code:**

```
1 from autogen import AssistantAgent, UserProxyAgent  
2  
3 assistant = AssistantAgent("assistant", llm="gpt-4")  
4 user = UserProxyAgent("user", code_execution=True)  
5  
6 user.initiate_chat(assistant, message="Build a weather app.")
```

OpenAI Agents

- ▶ **Type:** Hosted agent platform by OpenAI (2024+).
- ▶ **Features:**
 - ▶ Create, configure, and deploy GPT-based agents.
 - ▶ Access via OpenAI API and GPTs interface.
 - ▶ Integrates with files, APIs, tools, and memory.
 - ▶ Supports human feedback and evaluation.
- ▶ **Example Use:** Deploy a custom assistant with knowledge base and tools.

Google ADK (Agent Development Kit)

- ▶ **Developed by:** Google DeepMind / Gemini ecosystem.
- ▶ **Type:** Multi-agent and workflow orchestration framework.
- ▶ **Features:**
 - ▶ Integrates Gemini 1.5 models with tool-use capabilities.
 - ▶ Supports perception, memory, planning, and dialogue.
 - ▶ Human-in-the-loop and autonomous modes.
 - ▶ Built for scalable, production-grade AI agents.
- ▶ **Use Case:** Building Gemini-powered assistants and reasoning systems.

Comparison Summary

- ▶ **LangChain v1:** Graph-based, modular, integrates with LangSmith.
- ▶ **Pydantic-AI:** Strong typing, data-safe orchestration.
- ▶ **Crew AI:** Role-based multi-agent collaboration.
- ▶ **Autogen:** Conversational multi-agent system by Microsoft.
- ▶ **OpenAI Agents:** Hosted GPT-based no-code agent deployment.
- ▶ **Google ADK:** Gemini-native multi-agent orchestration suite.

Final Thoughts

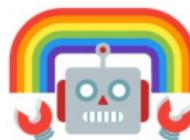
YHK

Why Agents?



Flexible

Can **reason** based on personalized inputs, handle unexpected **edge cases**, and respond in **natural language**. **Outcome-driven** rather than path-driven.



Easy to Prototype

Reduced dev time **integrating** with external APIs and databases, and in building deterministic, “if this then that” systems.



Proactive

Can run an agent in response to a **user prompt**, on a **trigger** or **schedule** (autonomously).

(Ref: Google Cloud Labs H2 India Agents for All)

When to Use Agents?

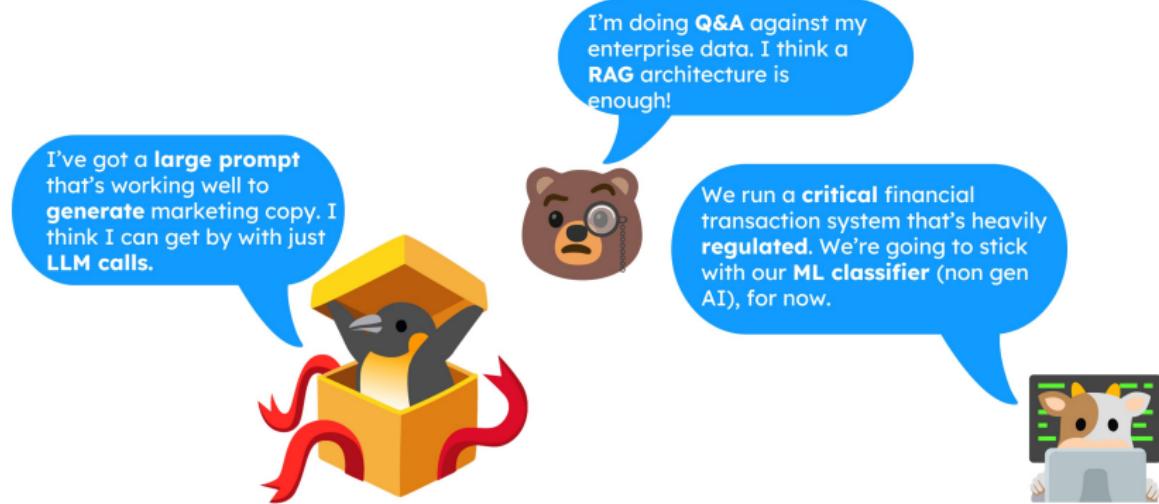
- ▶ Best suited for tasks requiring flexibility and model-driven decision-making
- ▶ Consider tradeoffs: agents increase latency and cost for better task performance
- ▶ Recommended for open-ended problems with unpredictable steps
- ▶ Simple solutions preferred - single LLM calls with retrieval often sufficient



(Ref: Google Cloud Labs H2 India Agents for All)

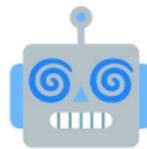
YHK

When NOT to Use Agents?



(Ref: Google Cloud Labs H2 India Agents for All)

Agent Challenges



Predictability

Agents use LLMs to reason and respond.
LLMs are nondeterministic.



Stability

Complexity of using many tools, **error-handling, security, privacy...**

+ Framework + protocol landscape is **evolving rapidly**.



Operations

Tracking calls through a **distributed system** is hard, debugging is tricky, tracking **token throughput** is important (**costs**), **quota**, LLM reasoning can be **opaque**...

(Ref: Google Cloud Labs H2 India Agents for All)

Mitigating Predictability challenges

- ▶ choose reasoning models, when possible. (eg. gemini 2.5 flash)
- ▶ ground LLM in trusted data to reduce hallucinations (RAG, search)
- ▶ guide your agent to think step by step (chain of thought or few-shot prompting)
- ▶ use state management and memory.
- ▶ implement checks in the agent's workflow: tool-call request validation, safety filters, logic checks...
- ▶ test and evaluate your agent.

(Ref: Google Cloud Labgs H2 India Agents for All)



Mitigating Stability challenges

- ▶ use consistent tool abstractions via Model Context Protocol (MCP)
- ▶ when writing your agent's system prompt, be as detailed as possible. (tool descriptions)
- ▶ implement or leverage error-handling in your agent (circuit-breakers, retries, hand off to a human in the loop...)
- ▶ utilize state management, especially the workflow status field (recover from interruptions, prevent infinite loops)
- ▶ be willing to refactor and adjust with new framework features, tool interfaces...

(Ref: Google Cloud Labs H2 India Agents for All)

Mitigating Operations challenges

- ▶ in the dev phase, log verbosely.
- ▶ build robust test and evaluation datasets. automate these tests on every commit.
- ▶ gather necessary metrics, like token throughput to your agent's models, and response code distribution. create alerts if these trip certain thresholds.
- ▶ leverage agent framework's tracing features to track model and tool calls, find + address latency bottlenecks.

(Ref: Google Cloud Labgs H2 India Agents for All)



Future AI Applications

- ▶ What are future AI applications like?
 - ▶ **Generative:** Generate content like text and images
 - ▶ **Agentic:** Execute complex tasks on behalf of humans
- ▶ How do we empower every developer to build them?
 - ▶ **Co-Pilots:** Human-AI collaboration
 - ▶ **Autonomous:** Independent task execution
- ▶ 2024 is expected to be the year of AI agents

The Big Question

- ▶ Agentic AI technology is moving incredibly fast
- ▶ Personal AI assistants will soon be capable of complex tasks
- ▶ Think beyond simple queries to multi-step accomplishments
- ▶ Consider what “impossible” tasks you want to delegate
- ▶ What complex, party-of-the-century level challenge will you tackle?
- ▶ The era of AI that truly does rather than just discusses
- ▶ Prepare for AI assistants that can handle your biggest challenges

Introduction to Multi Agent Systems

YHK

From Single to Multi-Agent Systems

- ▶ The Evolution of AI Systems:
 - ▶ QnA (Atomic Prompts) → Chatbot (Context)
 - ▶ RAG (Own Data) → Agents (Actions/Tools)
 - ▶ **Multi-Agent OS** (Orchestrated Intelligence)
- ▶ Multi-agent systems represent the pinnacle of AI orchestration
- ▶ Specialized AI agents collaborate to solve complex problems
- ▶ Each agent masters one framework deeply
- ▶ Together they achieve what no single expert could produce alone

Why Multiple Agents Beat Single Intelligence

- ▶ **Cognitive Bias:** Diverse agents challenge assumptions vs single perspective
- ▶ **Bounded Rationality:** Each agent specializes, synthesis integrates
- ▶ **Expertise Breadth:** Deploy 25+ specialist agents simultaneously
- ▶ **Consistency:** Agents maintain analytical rigor across perspectives
- ▶ **Parallel Processing:** Simultaneous analysis from multiple frameworks
- ▶ **Emotional Detachment:** Pure evaluation on merits without personal stakes
- ▶ **Scenario Analysis:** Each agent explores different timelines/strategies

The Wisdom of Artificial Crowds

- ▶ Based on **Society of Mind** (Marvin Minsky)
- ▶ Intelligence emerges from multiple simple processes
- ▶ Strategic wisdom emerges from:
 - ▶ **Specialized Analysis:** Each agent masters one framework
 - ▶ **Diverse Perspectives:** Game theorist vs mental model thinker
 - ▶ **Synthesis:** Oracle routes intelligently, Synthesizer integrates
 - ▶ **Emergent Insight:** The whole exceeds sum of parts
- ▶ Example: Business strategy analyzed by Nash, Prisoner's Dilemma, Coalition, and Mental Model agents
- ▶ Result: Strategic recommendation no single expert could produce alone

Multi-Agent Systems: Background

YHK

History of Multi-Agent Systems

- ▶ **1950s-1970s:** Early AI research on distributed problem solving
- ▶ **1980s:** Distributed Artificial Intelligence (DAI) emerges
 - ▶ Focus on coordinating multiple AI systems
 - ▶ Research in distributed reasoning and planning
- ▶ **1990s:** Multi-Agent Systems (MAS) becomes distinct field
 - ▶ Agent communication languages (KQML, FIPA-ACL)
 - ▶ Belief-Desire-Intention (BDI) architectures
- ▶ **2000s-2010s:** Practical applications emerge
 - ▶ Robotic swarms, game AI, trading systems
- ▶ **2020s:** LLM revolution transforms multi-agent systems
 - ▶ Natural language coordination becomes possible
 - ▶ Agents can reason, plan, and communicate effectively

Why Multi-Agent Systems?

Limitations of Single Agent Systems:

- ▶ Limited expertise scope - cannot excel at everything
- ▶ Single point of failure - no redundancy
- ▶ Bounded cognitive capacity - context window limits
- ▶ Lack of diverse perspectives - prone to bias

Advantages of Multi-Agent Systems:

- ▶ **Specialization:** Each agent masters specific domain
- ▶ **Parallelization:** Simultaneous processing of subtasks
- ▶ **Robustness:** System continues if one agent fails
- ▶ **Scalability:** Add agents for new capabilities
- ▶ **Modularity:** Easy to update or replace individual agents

Multi-Agent System Architectures: Overview

Three Main Architectural Patterns:

1. Hierarchical/Manager-Worker

- ▶ Central coordinator delegates to specialist agents
- ▶ Clear authority structure, predictable flow

2. Peer-to-Peer/Collaborative

- ▶ Agents communicate directly, negotiate solutions
- ▶ Democratic decision-making, emergent behavior

3. Pipeline/Sequential

- ▶ Each agent performs stage in workflow
- ▶ Output of one becomes input of next

Architecture 1: Hierarchical (Manager-Worker)

Structure:

- ▶ Manager agent coordinates
- ▶ Worker agents specialize
- ▶ Clear task delegation
- ▶ Centralized synthesis

Pros:

- ▶ Clear control flow
- ▶ Easy to debug
- ▶ Predictable behavior
- ▶ Prevents agent conflicts

Examples:

- ▶ AutoGen's group chat with manager
- ▶ CrewAI hierarchical mode
- ▶ MADS Oracle-Agent-Synthesizer

Cons:

- ▶ Manager bottleneck
- ▶ Less emergent behavior
- ▶ Single point of failure
- ▶ Limited agent autonomy

Architecture 2: Peer-to-Peer (Collaborative)

Structure:

- ▶ Agents communicate directly
- ▶ Democratic coordination
- ▶ Emergent solutions
- ▶ Consensus building

Pros:

- ▶ High agent autonomy
- ▶ Creative solutions
- ▶ No bottlenecks
- ▶ Robust to failures

Examples:

- ▶ AutoGen's autonomous chat
- ▶ MetaGPT role-playing
- ▶ ChatDev software team

Cons:

- ▶ Unpredictable behavior
- ▶ Hard to control
- ▶ May not converge
- ▶ Higher token costs

Architecture 3: Pipeline (Sequential)

Structure:

- ▶ Linear workflow stages
- ▶ Each agent processes then passes
- ▶ Assembly line pattern
- ▶ Clear input/output contracts

Pros:

- ▶ Simple to understand
- ▶ Easy to optimize stages
- ▶ Predictable flow
- ▶ Low coordination cost

Examples:

- ▶ LangGraph sequential chains
- ▶ Content creation pipelines
- ▶ Data processing workflows

Cons:

- ▶ Inflexible
- ▶ No parallel processing
- ▶ Errors cascade
- ▶ Limited adaptability

Communication Patterns in Multi-Agent Systems

- ▶ **Shared Memory/Blackboard**
 - ▶ Agents read/write to common state
 - ▶ LangGraph's state management
 - ▶ Good for: Coordinated information sharing
- ▶ **Message Passing**
 - ▶ Direct agent-to-agent communication
 - ▶ AutoGen's conversational approach
 - ▶ Good for: Negotiation, debate, collaboration
- ▶ **Publish-Subscribe**
 - ▶ Agents subscribe to topics of interest
 - ▶ Event-driven coordination
 - ▶ Good for: Decoupled, scalable systems

Coordination Mechanisms

How agents coordinate their actions:

- ▶ **Centralized Planning**

- ▶ Manager creates overall plan, assigns tasks
- ▶ Example: MADS Oracle routing to specialists

- ▶ **Distributed Planning**

- ▶ Each agent plans independently, negotiates conflicts
- ▶ Example: Market-based task allocation

- ▶ **Reactive Coordination**

- ▶ Agents respond to environment and each other
- ▶ Example: Swarm robotics, stigmergy

- ▶ **Contract Net Protocol**

- ▶ Agents bid for tasks, coordinator assigns to best bidder
- ▶ Example: Task allocation in heterogeneous systems

Multi-Agent Frameworks Comparison

AutoGen (Microsoft):

- ▶ Conversational agents
- ▶ Group chat patterns
- ▶ Code execution support
- ▶ Good for: Iterative tasks

CrewAI:

- ▶ Role-based agents
- ▶ Sequential/hierarchical modes
- ▶ Simple configuration
- ▶ Good for: Business workflows

LangGraph:

- ▶ Graph-based workflows
- ▶ Conditional routing
- ▶ State management
- ▶ Good for: Complex orchestration

MetaGPT:

- ▶ Software team simulation
- ▶ Standardized outputs (SOP)
- ▶ Document-driven
- ▶ Good for: Software development



Challenges in Multi-Agent Systems

► Coordination Overhead

- More agents = more communication = higher latency and cost
- Need efficient routing and synthesis mechanisms

► Consistency and Conflict Resolution

- Agents may produce contradictory recommendations
- Need robust conflict resolution and voting mechanisms

► Non-Determinism

- LLM variability multiplied across agents
- Difficult to guarantee reproducible results

► Debugging and Observability

- Hard to trace decisions through multi-agent interactions
- Need comprehensive logging and visualization tools

Best Practices for Multi-Agent Systems

- ▶ **Start Simple:** Begin with 2-3 agents, scale gradually
- ▶ **Clear Roles:** Define distinct, non-overlapping agent responsibilities
- ▶ **Explicit Communication:** Use structured formats (JSON, XML)
- ▶ **State Management:** Maintain shared context effectively
- ▶ **Termination Conditions:** Prevent infinite loops with clear stopping criteria
- ▶ **Human-in-Loop:** Include human oversight for critical decisions
- ▶ **Evaluation Metrics:** Define success criteria for system performance
- ▶ **Error Handling:** Graceful degradation when agents fail

Multi Agent Decision Systems (MADS)

YHK

MADS Core Philosophy

► Configuration Over Code

- Zero code required to build sophisticated systems
- Define everything in JSON configuration
- Domain-agnostic framework

► Multiple Domain Support

- Game Theory: 25+ agents for strategic decisions
- Mental Models: 10+ agents for personal decisions
- TRIZ Innovation: 40 inventive principles
- Business Strategy: Porter's Forces, Blue Ocean, etc.
- Custom Domains: Your expertise, your agents

Three-Stage Intelligence Architecture

Oracle (Classifier) → Specialist Agents (Analyzers) → Synthesizer (Integrator)

1. **Oracle Agent:** Analyzes problem, routes to 3-7 most relevant specialists
2. **Specialist Agents:** Each applies unique framework in parallel
3. **Synthesizer Agent:** Evaluates, ranks, integrates into unified recommendation

This three-agent core enables:

- ▶ Intelligent problem classification
- ▶ Parallel expert analysis
- ▶ Coherent strategy synthesis

How Problems Flow Through MADS

1. **Intake:** User describes strategic situation in natural language
2. **Classification:** Oracle extracts key dimensions (players, timing, uncertainty)
3. **Routing:** Oracle selects most relevant 3-7 specialist agents
4. **Parallel Analysis:** Selected agents simultaneously analyze from unique frameworks
5. **Evaluation:** Synthesizer scores each output on multiple criteria
6. **Integration:** Synthesizer resolves conflicts, finds synergies
7. **Output:** Executive summary + strategic guidance + implementation plan

Use Cases: Real-World Applications

YHK

Company-Level Strategic Decisions

Problem: "Should we acquire Competitor X or build feature Y internally?"

MADS Analysis:

- ▶ Game Theory agents evaluate competitive dynamics
- ▶ Coalition agent assesses M&A value distribution
- ▶ Opportunity cost agent quantifies what you give up
- ▶ Second-order thinking agent maps long-term consequences

Result: Nuanced recommendation with phased approach and decision triggers

Nation-Level Policy Analysis

Problem: "How should we respond to trade partner imposing tariffs?"

MADS Analysis:

- ▶ Repeated game agent models long-term diplomatic dynamics
- ▶ Nash equilibrium agent finds stable response strategies
- ▶ Bargaining agent identifies negotiation leverage
- ▶ Risk assessment agent evaluates escalation scenarios

Result: Multi-track strategy with clear red lines and off-ramps

Product Innovation

Problem: "Design a better smartphone charging solution"

MADS Analysis:

- ▶ TRIZ agents apply 40 inventive principles
- ▶ Systems thinking agent maps user journey and pain points
- ▶ First principles agent questions "why charging at all?"

Result: Portfolio of innovations from incremental to radical

Personal Life Decisions

Problem: "Should I take risky startup job or stay in stable corporate role?"

MADS Analysis:

- ▶ Opportunity cost agent quantifies what each path sacrifices
- ▶ Inversion agent identifies how each could fail
- ▶ Second-order consequences agent maps ripple effects
- ▶ Margin of safety agent recommends de-risking strategies

Result: Clear decision framework with contingency plans

Domain Frameworks

YHK

Game Theory Frameworks (25+ Agents)

Cooperative Game Theory

- ▶ Coalition Formation: Shapley values, core stability, optimal alliances
- ▶ Bargaining: Nash bargaining solution, BATNA analysis

Non-Cooperative Game Theory

- ▶ Nash Equilibrium: Pure/mixed strategies, stability analysis
- ▶ Dominant Strategy: Mechanism design

Classic Scenarios

- ▶ Prisoner's Dilemma, Stag Hunt, Chicken, Battle of Sexes

Mental Models (10+ Agents)

Cognitive frameworks for clearer thinking:

- ▶ **First Principles:** Break down to fundamental truths
- ▶ **Second-Order Thinking:** Consequences of consequences
- ▶ **Inversion:** Work backwards from failure to avoid it
- ▶ **Opportunity Cost:** What you sacrifice by choosing
- ▶ **Margin of Safety:** Build buffers for uncertainty
- ▶ **Circle of Competence:** Know what you know
- ▶ **Compound Interest:** How small advantages accumulate
- ▶ **Feedback Loops:** How systems reinforce or balance

TRIZ Innovation Principles (40 Agents)

Systematic innovation from studying 200,000+ patents:

- ▶ Segmentation, Asymmetry, Dynamization
- ▶ Preliminary Action, Cushion in Advance
- ▶ Self-Service, Copying, Cheap Disposables
- ▶ Nested Doll, Pneumatics, Flexible Shells
- ▶ Principle of universality, Mechanical vibration
- ▶ Thermal expansion, Strong oxidants
- ▶ And 25+ more systematic innovation principles

Each principle provides structured approach to innovation challenges

Configuration & Implementation

YHK

Creating Custom Configurations

Step 1: Define Your Domain Experts

- ▶ Identify specialist perspectives needed
- ▶ Example: Business Strategy domain
 - ▶ Porter's Five Forces Analyst
 - ▶ Blue Ocean Strategist
 - ▶ Disruptive Innovation Expert
 - ▶ Core Competency Evaluator

Step 2: Configure Agents in JSON

- ▶ Define agent ID, name, expertise areas
- ▶ Specify analysis framework and steps
- ▶ Set output format and prompt template

Oracle Routing & Synthesizer Evaluation

Step 3: Define Oracle Routing Logic

- ▶ Set conditions based on keywords or problem features
- ▶ Map conditions to relevant agent selections
- ▶ Ensure 3-7 agents selected for optimal analysis

Step 4: Configure Synthesizer Evaluation

- ▶ Define evaluation criteria (relevance, depth, feasibility)
- ▶ Assign weights to each criterion
- ▶ Common criteria: strategic clarity, competitive insight, actionability, risk awareness

Why LangGraph?

After evaluating AutoGen, CrewAI, and others:

- ▶ **Open Source:** Fully open with no vendor lock-in
- ▶ **Workflow Definition:** Excellent graph-based control
- ▶ **Open-Source LLMs:** Native support for local models
- ▶ **Conditional Routing:** Built-in intelligent routing
- ▶ **State Management:** Sophisticated state handling
- ▶ **Production Ready:** Deploy via LangServe
- ▶ **Community:** Large and active ecosystem

Decision: LangGraph for workflow control + open-source LLMs + production deployment



The Lang* Ecosystem

- ▶ **LangChain**: Framework for LLM applications, prompt chaining
- ▶ **LangGraph**: Multi-agent workflows with conditional routing (MADS uses this)
- ▶ **LangFlow**: Visual drag-and-drop interface for prototyping
- ▶ **LangSmith**: Testing, monitoring, debugging LLM applications
- ▶ **LangServe**: Production deployment with scaling and monitoring

Complete ecosystem for building, testing, and deploying agentic systems

Advanced Applications

Real-World MADS Applications

- ▶ **Quality Assurance:** Test-case generators from requirements
 - ▶ Edge cases, happy paths, error conditions, security threats
- ▶ **Brainstorming & Innovation:** Generate diverse ideas
 - ▶ TRIZ agents propose innovations, mental models challenge assumptions
- ▶ **Negotiation Optimization:** Multi-agent deliberation
 - ▶ Bargaining agents model BATNA, game theory finds equilibria
- ▶ **Investment Advisory:** Personalized portfolio recommendations
 - ▶ Risk profile, asset allocation, tax optimization, behavioral analysis

Current Limitations

Not Mission-Critical Ready Yet

- ▶ Non-deterministic outputs due to LLM variability
- ▶ Requires human-in-loop validation
- ▶ Best used in **Assistive/Co-pilot mode**, not full autonomy

Technical Limitations

- ▶ JSON parsing can fail with weaker LLMs
- ▶ Processing time: 30 seconds to 5 minutes per analysis
- ▶ Context window limits for very complex problems
- ▶ Rate limiting considerations for API-based LLMs



Roadmap to Production

Phase 1: Reliability (Q1 2026)

- ▶ Structured outputs enforcement, validation layers

Phase 2: Scale (Q2 2026)

- ▶ Parallel agent execution, caching, streaming outputs

Phase 3: Intelligence (Q3 2026)

- ▶ Agent learning from feedback, meta-agent improvements

Phase 4: Ecosystem (Q4 2026)

- ▶ Visual configuration builder, agent marketplace, enterprise features



Final Thoughts

YHK

The Big Picture: Multi-Agent Future

- ▶ **Vertical AI Agents could be 10X bigger than SaaS - Y Combinator**
- ▶ Moving from conversational AI to orchestrated intelligence
- ▶ Configuration-driven development: zero code for complex systems
- ▶ Domain-agnostic framework applicable to any field
- ▶ Open-source foundation with production-ready path
- ▶ 2025: The year of AI agents becomes reality
- ▶ Multi-agent systems solve problems humans struggle with alone
- ▶ We're not looking for Success but for **Wonder**



Key Takeaways

- ▶ Single agents are powerful; multi-agent systems are transformative
- ▶ Wisdom emerges from diverse perspectives working together
- ▶ Configuration over code enables rapid domain adaptation
- ▶ LangGraph provides production-ready orchestration
- ▶ Start simple: 3-7 agents for specific problems
- ▶ Scale complex: 25+ agents for strategic challenges
- ▶ Human-in-loop remains essential for reliability
- ▶ The future is collaborative intelligence, not singular AI

Next Steps

For Developers:

- ▶ Explore LangGraph and multi-agent frameworks
- ▶ Start with pre-built configurations (Game Theory, Mental Models)
- ▶ Create custom agents for your domain expertise

For Organizations:

- ▶ Identify high-value strategic decision points
- ▶ Pilot MADS on non-critical problems first
- ▶ Build domain-specific agent configurations
- ▶ Establish human oversight workflows

For Researchers:

- ▶ Contribute novel frameworks and evaluation metrics
- ▶ Improve agent coordination and synthesis algorithms

References

- ▶ CS 194/294-196 (LLM Agents) - Lecture 3, Chi Wang and Jerry Liu
- ▶ LLM Powered Autonomous Agents Lil'Log
- ▶ Power of Autonomous AI Agents - Yogesh Kulkarni
- ▶ Microsoft AutoGen- Yogesh Kulkarni
- ▶ Microsoft AutoGen using Open Source Models- Yogesh Kulkarni
- ▶ A CAMEL ride - Yogesh Kulkarni
- ▶ Autonomous AI Agents (LLM, VLM, VLA) - Code Your Own AI
- ▶ Awesome LLM-Powered Agent
<https://github.com/hyp1231/awesome-lilm-powered-agent>
- ▶ Autonomous Agents (LLMs). Updated daily
<https://github.com/tmgthb/Autonomous-Agents>

Thanks ...

- ▶ Office Hours: Saturdays, 3 to 5 pm (IST);
Free-Open to all; email for appointment to
yogeshkulkarni at yahoo dot com
- ▶ Call + 9 1 9 8 9 0 2 5 1 4 0 6



(<https://www.linkedin.com/in/yogeshkulkarni/>)



(<https://medium.com/@yogeshharibhaukulakarni>)



(<https://www.github.com/yogeshhk/>)



Pune AI Community (PAIC)

- ▶ Two-way communication:
 - ▶ Website puneaicommunity dot org
 - ▶ Email puneaicommunity at gmail dot com
 - ▶ Call + 9 1 9 8 9 0 2 5 1 4 0 6
 - ▶ LinkedIn:
<https://linkedin.com/company/pune-ai-community>
- ▶ One-way Announcements:
 - ▶ Twitter (X) @puneaicommunity
 - ▶ Instagram @puneaicommunity
 - ▶ WhatsApp Community: Invitation Link
<https://chat.whatsapp.com/LluOrhyEzuQLDr25ixZ>
 - ▶ Luma Event Calendar: puneaicommunity
- ▶ Contribution Channels:
 - ▶ GitHub: Pune-AI-Community and puneaicommunity
 - ▶ Medium: pune-ai-community
 - ▶ YouTube: @puneaicommunity



Website

Pune AI Community (PAIC) QR codes



Website



Medium Blogs



Twitter-X



LinkedIn Page



Github Repository



WhatsApp Invite



Luma Events



YouTube Videos



Instagram

